

CS101 Algorithms and Data Structures

Graphs

Textbook Ch B.4, B.5.1, 22.1

Outline

- Definitions
 - Undirected graphs
 - Directed graph
- Representation
 - Adjacency matrix
 - Adjacency list

Outline

A graph is an abstract data type for storing adjacency relations

- We start with definitions:
 - Vertices, edges, degree and sub-graphs
- We will describe paths in graphs
 - Simple paths and cycles
- Definition of connectedness
- Weighted graphs
- We will then reinterpret these in terms of directed graphs
- Directed acyclic graphs

Undirected Graphs

We will define an Undirected Graph ADT as a collection of *vertices*

$$V = \{v_1, v_2, \dots, v_n\}$$

- The number of vertices is denoted by

$$|V| = n$$


- Associated with this is a collection E of unordered pairs $\{v_i, v_j\}$ termed *edges* which connect the vertices

Undirected Graphs

Consider this collection of vertices

$$V = \{v_1, v_2, \dots, v_9\}$$

where $|V| = n = 9$




Undirected graphs

Associated with these vertices are $|E| = 5$ edges

$$E = \{\{v_1, v_2\}, \{v_3, v_5\}, \{v_4, v_8\}, \{v_4, v_9\}, \{v_6, v_9\}\}$$

- The pair $\{v_j, v_k\}$ indicates that both vertex v_j is adjacent to vertex v_k and vertex v_k is adjacent to vertex v_j



Undirected graphs

We will assume that a vertex is never adjacent to itself

- For example, $\{v_1, v_1\}$ will not define an edge

The maximum number of edges in an undirected graph is

$$|E| \leq \binom{|V|}{2} = \frac{|V|(|V|-1)}{2} = O(|V|^2)$$


An undirected graph

Example: given the $|V| = 7$ vertices

$$V = \{A, B, C, D, E, F, G\}$$

and the $|E| = 9$ edges

$$E = \{\{A, B\}, \{A, D\}, \{A, E\}, \{B, C\}, \{B, D\}, \{B, E\}, \{C, E\}, \{C, F\}, \{D, E\}\}$$



Degree


The degree of a vertex is defined as the number of adjacent vertices

$$\text{degree}(A) = \text{degree}(D) = \text{degree}(C) = 3$$

$$\text{degree}(B) = \text{degree}(E) = 4$$

$$\text{degree}(F) = 1$$



$$\text{degree}(G) = 0$$



Those vertices adjacent to a given vertex are its *neighbors*



Sub-graphs

A *sub-graph* of a graph contains a **subset** of the vertices and a subset of the edges that connect the **subset** of the vertices in the original graph





Sub-graphs

A *sub-graph* of a graph contains a subset of the vertices and a subset of the edges that connect the subset of the vertices in the original graph



Vertex-induced sub-graphs

A *vertex-induced sub-graph* contains a subset of the vertices and all the edges in the original graph between those vertices



Paths

A path in an undirected graph is an ordered sequence of vertices

$(v_0, v_1, v_2, \dots, v_k)$


where $\{v_{j-1}, v_j\}$ is an edge for $j = 1, \dots, k$

- Termed *a path from v_0 to v_k*
- The length of this path is k

Paths

A path of length 4:


(A, B, E, C, F)



Paths

A path of length 5:


(A, B, E, C, B, D)



Paths


A trivial path of length 0:

(A)




Simple path

A *simple path* has no repetitions (other than perhaps the first and last vertices)




(A, B, E, C, F)



(A, B, E, C, B, D)

Simple cycle

A *simple cycle* is a simple path of **at least two vertices** with the first and last vertices equal




(A, B, D, A)


Connectedness

Two vertices v_i, v_j are said to be *connected* if there exists a path from v_i to v_j

A graph is connected if there exists a path between any two vertices



A connected graph




An unconnected graph

Weighted graphs

A weight may be associated with each edge in a graph


- This could represent distance, energy consumption, cost, etc.
- Such a graph is called a *weighted graph*



Weighted graphs

The *length* of a path within a weighted graph is the sum of all of the edges which make up the path


- The length of the path (A, D, G) in the following graph is $5.1 + 3.7 = 8.8$



Weighted graphs

Different paths may have different weights


- Another path is (A, C, F, G) with length $1.2 + 1.4 + 4.5 = 7.1$



Weighted graphs

Problem: find the shortest path between two vertices




- Here, the shortest path from A to G is (A, C, F, D, E, G) with length 5.7



Trees

A graph is a tree if it is connected and there is a unique path between any two vertices

- Example: three trees on the same eight vertices



Properties:

- The number of edges is $|E| = |V| - 1$
- The graph is *acyclic*, that is, it does not contain any cycles
- Adding one more edge must create a cycle
- Removing any one edge creates two unconnected sub-graphs


Trees

Any tree can be converted into a rooted tree by:

- Choosing any vertex to be the root
- Defining its neighboring vertices as its children

and then recursively defining:

- All neighboring vertices other than that one designated its parent to be its children



Forests


A forest is any graph that has no cycles

Consequences:

- The number of edges is $|E| < |V|$
- The number of trees is $|V| - |E|$
- Removing any one edge adds one more tree to the forest

Here is a forest with 22 vertices and 18 edges

- There are four trees



Outline

- Definitions
 - Undirected graphs
 - Directed graph
- Representation
 - Adjacency matrix
 - Adjacency list

Directed graphs

In a *directed graph*, the **edges** on a graph are be **associated with a direction**

- Edges are ordered pairs (v_j, v_k) denoting a connection from v_j to v_k
- The edge (v_j, v_k) is different from the edge (v_k, v_j)

Streets are directed graphs:


- In most cases, you can go two ways unless it is a one-way street

Directed graphs

Given a graph of nine vertices $V = \{v_1, v_2, \dots, v_9\}$

- These six pairs (v_j, v_k) are *directed edges*

$$E = \{(v_1, v_2), (v_3, v_5), (v_5, v_3), (v_6, v_9), (v_8, v_4), (v_9, v_4)\}$$



Directed graphs

The maximum number of directed edges in a directed graph is

$$|E| \leq 2 \binom{|V|}{2} = 2 \frac{|V|(|V|-1)}{2} = |V|(|V|-1) = O(|V|^2)$$

In and out degrees


The degree of a vertex in a directed graph:

- The *out-degree* of a vertex is the number of outward edges from the vertex
- The *in-degree* of a vertex is the number of inward edges to the vertex

In this graph:

$$\text{in_degree}(v_1) = 0 \quad \text{out_degree}(v_1) = 2$$

$$\text{in_degree}(v_5) = 2 \quad \text{out_degree}(v_5) = 3$$




Sources and sinks

Definitions:

- Vertices with an in-degree of zero are described as *sources*
- Vertices with an out-degree of zero are described as *sinks*

In this graph:

- Sources: v_1, v_6, v_7
- Sinks: v_2, v_9



Paths

A path in a directed graph is an ordered sequence of vertices

$$(v_0, v_1, v_2, \dots, v_k)$$


where (v_{j-1}, v_j) is an edge for $j = 1, \dots, k$

A path of length 5 in this graph is

$$(v_1, v_4, v_5, v_3, v_5, v_2)$$

A simple cycle of length 3 is

$$(v_8, v_4, v_5, v_8)$$




Connectedness

Two vertices v_j, v_k are said to be *connected* if there exists a path from v_j to v_k

- A graph is *strongly connected* if there exists a directed path between any two vertices
- A graph is *weakly connected* if there exists a path between any two vertices that ignores the direction

In this graph:


- The sub-graph $\{v_3, v_4, v_5, v_8\}$ is strongly connected
- The sub-graph $\{v_1, v_2, v_3, v_4, v_5, v_8\}$ is weakly connected



Weighted directed graphs

In a weighted directed graphs, each edge is associated with a value

If both (v_j, v_k) and (v_k, v_j) are edges, it is not required that they have the same weight




Directed acyclic graphs


A *directed acyclic graph* is a directed graph which has **no cycle**

- These are commonly referred to as DAGs
- They are graphical representations of partial orders on a finite number of elements

These two are DAGs:



This directed graph is not acyclic:



Directed acyclic graphs

Applications of directed acyclic graphs include:


- The parse tree constructed by a compiler
- A reference graph that can be garbage collected using simple reference counting
- Dependency graphs such as those used in instruction scheduling and **makefiles**
- Dependency graphs between classes formed by inheritance relationships in object-oriented programming languages
- Information categorization systems, such as folders in a computer
- Directed acyclic word graph data structure to memory-efficiently store a set of strings (words)

Reference: http://en.wikipedia.org/wiki/Directed_acyclic_graph

Representations

How do we store the adjacency relations?

- Binary-relation list
- Adjacency matrix
- Adjacency list



Summary

In this topic, we have covered:

- Basic graph definitions
 - Vertex, edge, degree, adjacency
- Paths, simple paths, and cycles
- Connectedness
- Weighted graphs
- Directed graphs
- Directed acyclic graphs

We will continue by looking at a number of problems related to graphs

References

Wikipedia, http://en.wikipedia.org/wiki/Adjacency_matrix
http://en.wikipedia.org/wiki/Adjacency_list

- [1] Donald E. Knuth, *The Art of Computer Programming, Volume 1: Fundamental Algorithms*, 3rd Ed., Addison Wesley, 1997, §2.2.1, p.238.
- [2] Cormen, Leiserson, and Rivest, *Introduction to Algorithms*, McGraw Hill, 1990, §11.1, p.200.
- [3] Weiss, *Data Structures and Algorithm Analysis in C++*, 3rd Ed., Addison Wesley, §3.6, p.94.
- [4] David H. Laidlaw, Course Notes, <http://cs.brown.edu/courses/cs016/lectures/13%20Graphs.pdf>

These slides are provided for the ECE 250 *Algorithms and Data Structures* course. The material in it reflects Douglas W. Harder's best judgment in light of the information available to him at the time of preparation. Any reliance on these course slides by any party for any other purpose are the responsibility of such parties. Douglas W. Harder accepts no responsibility for damages, if any, suffered by any party as a result of decisions made or actions based on these course slides for any other purpose than that for which it was intended.

Outline

- Definitions
 - Undirected graphs
 - Directed graph
- Representation
 - Adjacency matrix
 - Adjacency list

The Graph ADT

The Graph ADT describes a container storing an adjacency relation

– Queries include:

- The number of vertices
- The number of edges
- List the vertices adjacent to a given vertex
- Are two vertices adjacent?
- Are two vertices connected?

– Modifications include:

- Inserting or removing an edge
- Inserting or removing a vertex (and all edges containing that vertex)

The run-time of these operations will depend on the representation


Binary-relation list

The most inefficient is a relation list:

- A container storing the edges

$\{(1, 2), (1, 4), (3, 5), (4, 2), (4, 5), (5, 2), (5, 3), (5, 8), (6, 9), (7, 9), (8, 4)\}$

- Requires $\Theta(|E|)$ memory
- Determining if v_j is adjacent to v_k is $O(|E|)$
- Finding all neighbors of v_j is $\Theta(|E|)$




Adjacency matrix

Requiring more memory but also faster, an adjacency matrix

- The matrix entry (j, k) is set to true if there is an edge (v_j, v_k)

	1	2	3	4	5	6	7	8	9
1		T		T					
2									
3					T				
4		T			T				
5		T	T					T	
6									T
7								T	
8									
9									

- Requires $\Theta(|V|^2)$ memory
- Determining if v_j is adjacent to v_k is $O(1)$
- Finding all neighbors of v_j is $\Theta(|V|)$



Adjacency list

Most efficient for algorithms is an adjacency list

- Each vertex is associated with a list of its neighbors

```

1  • → 2 → 4
2  •
3  • → 5
4  • → 2 → 5
5  • → 2 → 3 → 8
6  • → 9
7  • → 9
8  • → 4
9  •
    
```

- Requires $\Theta(|V| + |E|)$ memory


- On average:

- Determining if v_j is adjacent to v_k is

- Finding all neighbors of v_j is

$$O\left(\frac{|E|}{|V|}\right)$$

$$\cdot \left(\frac{|E|}{|V|}\right)$$



Outline

- In this topic, we will cover the representation of graphs on a computer
- We will examine:
 - an adjacency matrix representation
 - smaller representations and pointer arithmetic
 - sparse matrices and linked lists

Adjacency Matrix

A graph of n vertices may have up to

$$\binom{n}{2} = \frac{n(n-1)}{2} = \mathbf{O}(n^2)$$

edges


The first straight-forward implementation is an adjacency matrix

Adjacency Matrix

Define an $n \times n$ matrix $\mathbf{A} = (a_{ij})$ and if the vertices v_i and v_j are connected with weight w , then set $a_{ij} = w$ and $a_{ji} = w$

That is, the matrix is symmetric, e.g.,

	0	1	2	3	4	5
0		0.83		0.65		0.51
1	0.83		0.67	0.72	0.46	0.90
2		0.67		0.39		
3	0.65	0.72	0.39		0.24	
4		0.46		0.24		
5	0.51	0.90				




Adjacency Matrix

An unweighted graph may be saved as an array of Boolean values

- vertices v_i and v_j are connected then set

$$a_{ij} = a_{ji} = \text{true}$$


	0	1	2	3	4	5
0		T	F	T	F	T
1	T		T	T	T	T
2	F	T		T	F	F
3	T	T	T		T	F
4	F	T	F	T		F
5	T	T	F	F	F	



Adjacency Matrix

If the graph was directed, then the matrix would not necessarily be symmetric

	0	1	2	3	4	5
0				0.65		
1	0.83			0.72	0.46	0.90
2		0.67		0.39		
3					0.24	
4				0.24		
5	0.51					



Adjacency Matrix

First we must allocate memory for a two-dimensional array

C++ does not have native support for anything more than one-dimensional arrays, thus how do we store a two-dimensional array?

- as an array of arrays


Adjacency Matrix

Suppose we require a 16×16 matrix of double-precision floating-point numbers

Each row of the matrix can be represented by an array

The address of the first entry must be stored in a pointer to a double:

```
double *
```




Adjacency Matrix

However, because we must store 16 of these pointers-to-doubles, it makes sense that we store these in an array

What is the declaration of this array?

Well, we must store a
pointer to a pointer to a double


That is: `double **`



Adjacency Matrix

Thus, the address of the first array must be declared to be:

```
double **matrix;
```



Adjacency Matrix

The next question is memory allocation

First, we must allocate the memory for the array of pointers to doubles:


```
matrix = new double * [16];
```



Adjacency Matrix

Next, to each entry of this matrix, we must assign the memory allocated for an array of doubles

```
for ( int i = 0; i < 16; ++i ) {  
    matrix[i] = new double[16];  
}
```




Adjacency Matrix

Accessing a matrix is done through a double index, *e.g.*,
`matrix[3][4]`

You can interpret this as `(matrix[3])[4]`


Adjacency Matrix

Recall that in `matrix[3][4]`, the variable `matrix` is a pointer-to-a-pointer-to-a-double:




Adjacency Matrix

Therefore, `matrix[3]` is a pointer-to-a-double:



Adjacency Matrix

And consequently, `matrix[3][4]` is a double:



C++ Notation Warning

Do not use `matrix[3, 4]` because:

- in C++, the comma operator evaluates the operands in order from left-to-right
- the *value* is the last one

Therefore, `matrix[3, 4]` is equivalent to calling `matrix[4]`

Try it:

```
int i = (3, 4);  
cout << i << endl;
```

C++ Notation Warning

Many things will compile if you try to use this notation:

```
matrix = new double[N, N];
```

will allocate an array of N doubles, just like:

```
matrix = new double[N];
```


However, this is likely not to do what you really expect...

Adjacency Matrix

Recall that for each call to `new[]`, you must have a corresponding call to `delete[]`

Therefore, we must use a for-loop to delete the arrays

- implementation up to you



Default Values

Question: what do we do about vertices which are not connected?

- the value 0
- a negative number, e.g., -1
- positive infinity: ∞

The last is the most logical, in that it makes sense that two vertices which are not connected have an infinite distance between them

The distance from a node to itself is 0

Default Values

To use infinity, you may declare a constant static member variable INF:

```
#include <limits>
```

```
class Weighted_graph {  
    private:  
        static const double INF;  
        // ...  
        // ...  
};
```

```
const double Weighted_graph::INF =  
    std::numeric_limits<double>::infinity();
```

Default Values

As defined in the IEEE 754 standard, the representation of the double-precision floating-point infinity eight bytes:

0x 7F F0 00 00 00 00 00 00

Incidentally, negative infinity is stored as:

0x FF F0 00 00 00 00 00 00

Default Values

In this case, you can initialize your array as follows:

```
for ( int i = 0; i < N; ++i ) {  
    for ( int j = 0; j < N; ++j ) {  
        matrix[i][j] = INF;  
    }  
  
    matrix[i][i] = 0;  
}
```

It makes intuitive sense that the distance from a node to itself is 0

Default Values

If we are representing an unweighted graph, use Boolean values:

```
for ( int i = 0; i < N; ++i ) {  
    for ( int j = 0; j < N; ++j ) {  
        matrix[i][j] = false;  
    }  
  
    matrix[i][i] = true;  
}
```

It makes intuitive sense that a vertex is connected to itself

Sparse Matrices

- The memory required for creating an $n \times n$ matrix using a 2D array is $\Theta(n^2)$ bytes
- This could potentially waste a significant amount of memory:
 - Consider a friendship graph: nodes represent persons and edges represent friendship
 - The world population is 7.4 billion \Rightarrow the size of the matrix is $(7.4 \times 10^9)^2 \approx 55 \times 10^{18}$
 - However, each person on average has, say, 100 friends. Hence only $\frac{100}{7.4 \times 10^9}$ of the matrix elements are true. The other elements are the default value: false.

Sparse Matrices

- Matrices where less than 5% of the entries are not the default value (either infinity or 0, or perhaps some other default value) are said to be *sparse*
- Matrices where most entries (25% or more) are not the default value are said to be *dense*
- Clearly, these are not hard limits

Adjacency list

- For an undirected graph, use an array of linked lists to store edges
 - Each vertex has a linked list that stores all the edges connected to the vertex
 - Each node in a linked list must store two items of information: the connecting vertex and the weight

Adjacency list

We may create a new class which stores a vertex-edge pair

```
class Pair {  
    private:  
        double edge_weight;  
        int adjacent_vertex;  
    public:  
        Pair( int, double );  
        double weight() const;  
        int vertex() const;  
};
```


Now create an array of linked-lists storing these pairs

Adjacency list

Thus, we define and create the array:

```
SingleList<Pair> * array;
```

```
array = new SingleList<Pair>[16];
```




Adjacency list

To reduce redundancy, we would only insert the pair into the linked list corresponding to the larger vertex



```
void insert( int i, int j, double w ) {  
    if ( i < j ) {  
        array[j].push_front( Pair(i, w) );  
    } else {  
        array[i].push_front( Pair(j, w) );  
    }  
}
```

Adjacency list

For example, the graph shown below would be stored as





Adjacency list



Adjacency list

- To store a **directed graph**
 - Each vertex has a linked list that stores all the edges originated from the vertex
 - Each node in a linked list stores two items of information: the vertex that the edge connects to, the weight

Adjacency list



Summary

- In this laboratory, we have looked at a number of graph representations
- C++ lacks a *matrix* data structure
 - must use array of arrays
- The possible factors affecting your choice of data structure are:
 - weighted or unweighted graphs
 - directed or undirected graphs
 - dense or sparse graphs

References

Wikipedia, http://en.wikipedia.org/wiki/Adjacency_matrix
http://en.wikipedia.org/wiki/Adjacency_list

These slides are provided for the ECE 250 *Algorithms and Data Structures* course. The material in it reflects Douglas W. Harder's best judgment in light of the information available to him at the time of preparation. Any reliance on these course slides by any party for any other purpose are the responsibility of such parties. Douglas W. Harder accepts no responsibility for damages, if any, suffered by any party as a result of decisions made or actions based on these course slides for any other purpose than that for which it was intended.

Summary

- Definitions
 - Undirected graphs
 - Directed graph
- Representation
 - Adjacency matrix
 - Adjacency list