



信息科学与技术学院

School of Information Science and Technology

# CS 110

# Computer Architecture

# C Pointers & Array

**Instructors:**

**Siting Liu & Chundong Wang**

Course website: [https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/  
Spring-2023/index.html](https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html)

**School of Information Science and Technology (SIST)**

**ShanghaiTech University**


2023/2/6

# Outline

- Pointer
- Array
- Memory Management


# Address vs. Value

- Consider memory to be a single huge array
  - Each cell of the array has an address associated with it
  - Each cell also stores some value
  - For addresses do we use signed or unsigned numbers?
- Don't confuse the address referring to a memory location with the value stored there




# Pointers

- An address refers to a particular memory location; e.g., it points to a memory location
- Pointer: A variable that contains the address of a variable



# Pointer Syntax


- `int *p;`
  - Tells compiler that `variable p` is the `address` of an `int`
- `p = &X;`
  - Tells compiler to assign `address of X` to `p`
  - `&` called the “address operator” in this context
- `z = *p;`
  - Tells compiler to assign `value at address p` to `z`
  - `*` called the “dereference operator” in this context
- Can point to `int/char/struct/function`, etc.



# Creating and Using Pointers

- How to create a pointer:

& operator: get address of a variable



Note the “\*” gets used 2 different ways in this example. In the declaration to indicate that **p** is going to be a pointer, and in the **printf** to get the value pointed to by **p**.


- How to get a value pointed to?

“\*” (dereference operator): get the value that the pointer points to

```
printf("p points to value %d\n", *p);
```

# Using Pointer for Writes

- How to change a variable pointed to?
  - Use the dereference operator **\*** on the left of assignment operator **=**



# Pointers and Parameter Passing

- C passes parameters “by value”
  - Procedure/function/method gets a copy of the parameter, so *changing the copy cannot change the original*

```
void add_one (int x) {  
    x = x + 1;  
}  
  
int y = 3;  
add_one(y);
```

*y remains equal to 3*



# Pointers and Parameter Passing

- How can we get a function to change the value held in a variable?

```
void add_one (int *p) {  
    *p = *p + 1;  
}
```

```
int y = 3;
```

```
add_one(&y);
```

*y is now equal to 4*

# Types of Pointers

- Pointers are used to point to any kind of data (**int**, **char**, a **struct**, etc.)
- Normally a pointer only points to one type (**int**, **char**, a **struct**, etc.).
  - **void \*** is a type that can point to anything (generic pointer)
  - Be careful, use **void \*** sparingly to help avoid program bugs, and security issues, and other bad things!

# More C Pointer Dangers

- Declaring a pointer just allocates space to hold the pointer – it does not allocate the thing being pointed to!
- Local variables in C are not initialized, they may contain anything (a.k.a. “garbage”)
- What does the following code do?

Undefined Behavior

```
void f()  
{  
    int *ptr;  
    *ptr = 5;  
}
```

Not even working in  
some environments

# Pointers and Structures

```
typedef struct {                /* dot notation */
    int x;                      int h = p1.x;
    int y;                      p2.y = p1.y;
} Point;

                                /* arrow notation */
Point p1;                      int h = paddr->x;
Point p2;                      int h = (*paddr).x;
Point *paddr;

                                /* This works too */
p1 = {1,2};                    p1 = p2;
```

# Pointers in C

- Why use pointers?
  - If we want to pass a large struct or array, it's easier/faster/etc. to pass a pointer than the whole thing
  - In general, pointers allow cleaner, more compact code
- So what are the drawbacks?
  - Pointers are probably the single largest source of bugs in C, so be careful anytime you deal with them
    - Most problematic with dynamic memory management—  
[coming up soon](#)
    - Dangling references and memory leaks

# Why Pointers in C?

- At time C was invented (early 1970s), compilers often didn't produce efficient code
  - Computers 100,000 times faster today, compilers better
- C designed to let programmer say what they want code to do without compiler getting in way
  - Even give compilers hints which registers to use!
- Today's compilers produce much better code, so may not need to use pointers in application code
- Low-level system code still needs low-level access via pointers

# Course Info

- HW1 due today! HW2 will be available very soon (see piazza & course webpage). Lab 2 will be available this Friday.
- Find a lab-mate in the same lab session, two students in a group, and Sunday deadline approaching!
- Lab 1 is available and in this week's Lab session
- Lab & Discussion (see webpage, this Friday, VMware/Ubuntu linux, etc.) starts this week.
- **Webpage: <https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html>**
- Discussion next week: Venus (for RISC-V), Memory Management & debug

# Outline

- Pointer
- **Array**
- Memory Management



# C Arrays

- Declaration:

```
int ar[2];
```

declares a 2-element integer array: just a block of memory

```
int ar[] = {795, 635};
```

declares and initializes a 2-element integer array

- Must specify size (or provide info. that can infer the size)

# C Strings

- String in C is just an array of characters

```
char string[] = "abc";
```

- How do you tell how long a string is?
  - Last character is followed by a ‘\0’ (NULL) byte (a.k.a. “null terminator”) (RTFM)

```
int strlen(char s[])  
{  
    int n = 0;  
    while (s[n] != 0)  
        n++;  
    return n;  
}
```

Finally, the declaration

```
char s[] = "abc", t[3] = "abc";
```

defines ‘plain’ char array objects s and t whose members are initialized with character string literals. This declaration is identical to

```
char s[] = { 'a', 'b', 'c', '\0' },  
t[] = { 'a', 'b', 'c' };
```

# Array Name/Pointer Duality

- Key Concept: Array variable is a “pointer” to the first (lowest addressed, i.e., 0<sup>th</sup>) element
- So, array variables almost identical to pointers
  - **char \*string** and **char string[]** are nearly identical declarations
  - Differ in subtle ways: incrementing, declaration of filled arrays
- Consequences:
  - **ar** is an array variable, but works like a pointer
  - **ar[0]** is the same as **\*ar**
  - **ar[2]** is the same as **\*(ar+2)**
  - Can use pointer arithmetic to conveniently access arrays

# Array Name/Pointer Duality

- Be really careful!

```
char string1[] = "abc";
```

```
char *string2 = "abc";
```

- **CAN** modify string1[\*]
- **CANNOT** modify string2[\*]
- **CAN** access string2 by string2[\*]

# Array Name/Pointer Duality

- Be really careful!

```
int arr[] = { 3, 5, 6, 7, 9 };
```

```
int *p = arr;
```

```
int (*p1)[5] = &arr;
```

```
int *p2[5];
```

```
int *p3(void);
```


- Read More: C manual 3.5.5 Type names!

# Changing a Pointer Argument?

- What if want function to change a pointer?
- What gets printed?

```
void inc_ptr(int *p)
{
    p = p + 1;
}

int A[3] = {50, 60, 70};
int *q = A;
inc_ptr(q);
printf("q = %d\n", *q);
```



# Pointer to a Pointer


- Solution! Pass a pointer to a pointer, declared as **\*\*h**

- Now what gets printed?

Pointer to pointer to an int

```
void inc_ptr(int **h)
{
    *h = *h + 1;
}

int A[3] = {50, 60, 70};
int *q = A;
inc_ptr(&q);
printf("*q = %d\n", *q);
```



# C Arrays are Very Primitive

- An array in C does not know its own length, and its bounds are not checked!
  - Consequence: We can accidentally access off the end of an array
  - Suggestion: We must pass the array and its size to any procedure that is going to manipulate it
- Out of boundary errors:
  - These are VERY difficult to find;  
Be careful!



# Use Defined Constants

- Array size  $n$ ; want to access from 0 to  $n-1$ , so you should use counter AND utilize a variable for declaration & incrementation
  - Bad pattern

```
int i, ar[10];
for(i = 0; i < 10; i++){ ... }
```
  - Better pattern

```
const int ARRAY_SIZE = 10;
int i, a[ARRAY_SIZE];
for(i = 0; i < ARRAY_SIZE; i++){ ... }
```
- SINGLE SOURCE OF TRUTH
  - You avoid maintaining two copies of the number 10
  - DRY: “Don’t Repeat Yourself”

# Pointing to Different Size Objects

- Modern machines are “byte-addressable”
  - Hardware’s memory composed of 8-bit storage cells, each has a unique address
- A C pointer is just abstracted memory address
- Type declaration tells compiler how many bytes to fetch on each access through pointer
  - E.g., 32-bit integer stored in 4 consecutive 8-bit bytes


- Alignment

```
%1 = alloca i32, align 4
%2 = alloca i32, align 4
%3 = alloca i32, align 4
%4 = alloca i32, align 4
store i32 0, i32* %1, align 4
store i32 1234, i32* %2, align 4
store i32 4321, i32* %3, align 4
%5 = load i32, i32* %2, align 4
%6 = load i32, i32* %3, align 4
%7 = add nsw i32 %5, %6
store i32 %7, i32* %4, align 4
```

**char \*z**  
8-bit character  
stored in one byte

**short \*y**  
16-bit short stored  
in two bytes

**int \*x**  
32-bit integer  
stored in four bytes



# sizeof() operator

- sizeof(type) returns number of bytes in object
  - But number of bits in a byte is not standardized
    - In olden times, when dragons roamed the earth, bytes could be 5, 6, 7, 9 bits long
- By definition, sizeof(char)==1
- Can take sizeof(variable), or sizeof(type)
- We'll see more of sizeof when we look at dynamic memory management

# Pointer Arithmetic

*pointer + number*      *pointer – number*

e.g., *pointer + 1*      adds 1 to a pointer, used in array access

```
char    *p;  
char    a;  
char    b;  
  
p = &a;  
p += 1;
```

In each, p now points to b

(Assuming compiler doesn't  
reorder variables in memory.

**Never code like this!!!!)**

```
int      *p;  
int      a;  
int      b;  
  
p = &a;  
p += 1;
```

Adds **1\*sizeof(char)**  
to the memory address

Adds **1\*sizeof(int)**  
to the memory address

*Pointer arithmetic should be used cautiously*

# Arrays and Pointers

## Passing arrays

- Array  $\approx$  pointer to the initial (0<sup>th</sup>) array element  
 $\mathbf{a[i] \equiv (* (a + (i)))}$
- An array is passed to a function as a pointer
  - The array size is lost!
- Usually bad style to interchange arrays and pointers
  - Avoid pointer arithmetic!

Really `int *array`      Must explicitly pass the size

```
int
foo(int array[],
    unsigned int size)
{
    ... array[size - 1] ...
}

int
main(void)
{
    int a[10], b[5];
    ... foo(a, 10) ... foo(b, 5) ...
}
```

# Arrays and Pointers

```
int
foo(int array[],
    unsigned int size)
{
    ...
    printf("%d\n", sizeof(array));
}
```

Compiler time operator

```
int
main(void)
{
    int a[10], b[5];
    ... foo(a, 10)... foo(b, 5) ...
    printf("%d\n", sizeof(a));
}
```

What does this print  
(32-bit address)?

4

... because **array** is really  
a pointer (and a pointer is  
architecture dependent, but  
likely to be 8 on modern  
machines!)

What does this print  
(32-bit int)?

40

# Arrays and Pointers

```
int i;  
int array[10];  
  
for (i = 0; i < 10; i++)  
{  
    array[i] = ...;  
}
```

```
int *p;  
int array[10];  
  
for (p = array; p < &array[10]; p++)  
{  
    *p = ...;  
}
```

These code sequences have the same effect!

# Concise strlen()

```
int(long) strlen(char *s)
{
    char *p = s;
    while (*p++)
        ; /* Null body of while */
    return (p - s - 1);
}
```

What happens if there is no zero character at end of string?



# Point past end of array?

- Array size n; want to access from 0 to n-1, but test for exit by comparing to address one element past the last member of the array

```
int ar[10]={},*p, *q, sum=0;
p = &ar[0]; q = &ar[10];
while (p!=q) /* sum = sum+*p; p = p+1*/
    sum += *p++;
```

- C defines that one element past end of array must be a valid address, i.e., not causing an error

# Valid Pointer Arithmetic

- Add an integer to a pointer.
- Subtract 2 pointers (pointed to the same array/object)
- Compare pointers (<, <=, ==, !=, >, >=)
- Compare pointer to **NULL** (indicates that the pointer 0x0)

Everything else illegal since makes no sense:

- adding two pointers
- multiplying pointers
- subtract pointer from integer
- etc.

# Arguments in `main()`

- To get arguments to the main function, use:
  - `int main(int argc, char *argv[])`
- What does this mean?
  - `argc` contains the number of strings on the command line (the executable counts as one, plus one for each argument). Here `argc` is 5:  
`% clang -ansi introC_1_1.c -o introC_1_1.out`
  - `argv` is a pointer to an array containing the arguments as strings

# Example

```
% clang -ansi introC_1_1.c -o introC_1_1.out
```

- `argc = 5 /* number arguments */`
- `argv[0] = "clang",`  
`argv[1] = "-ansi",`  
`argv[2] = "introC_1_1.c",`  
`argv[3] = "-o",`  
`argv[4] = "introC_1_1.out",`
- Array of pointers to strings

# Summary

- “Lowest High-level language”
  - Use ANSI C89 in class
  - => closest to assembler
- Pointers: powerful but dangerous
- Pointer arithmetic and arrays useful but also dangerous

# Summary

- Pointers and arrays are **virtually same**
- C **knows** how to increment pointers
- C is an efficient language, with little protection
  - Array bounds **not checked**
  - Variables **not** automatically initialized
- (Beware) The cost of efficiency is more overhead for the programmer.

**“C gives you a lot of extra rope but be careful not to hang yourself with it!”**



信息科学与技术学院

School of Information Science and Technology

# CS 110

# Computer Architecture

# C Memory Management

**Instructors:**

**Siting Liu & Chundong Wang**

Course website: <https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html>

**School of Information Science and Technology (SIST)**

**ShanghaiTech University**

2023/2/6

# C Memory Management

- To simplify, assume one program runs at a time

- A program's address space contains 4 regions:

Memory Address  
(32 bits assumed here)

- stack: local variables inside functions, grows downward


$\sim FFFF\ FFFF_{hex}$

- heap: space requested for dynamic data via `malloc()`; resizes dynamically, grows upward

- static {
- static data: variables declared outside functions, does not grow or shrink. Loaded when program starts, can be modified.

- code (a.k.a. text): loaded when program starts, does not change

- 0x0 unwritable/unreadable (NULL pointer)





# Where are Variables Allocated?

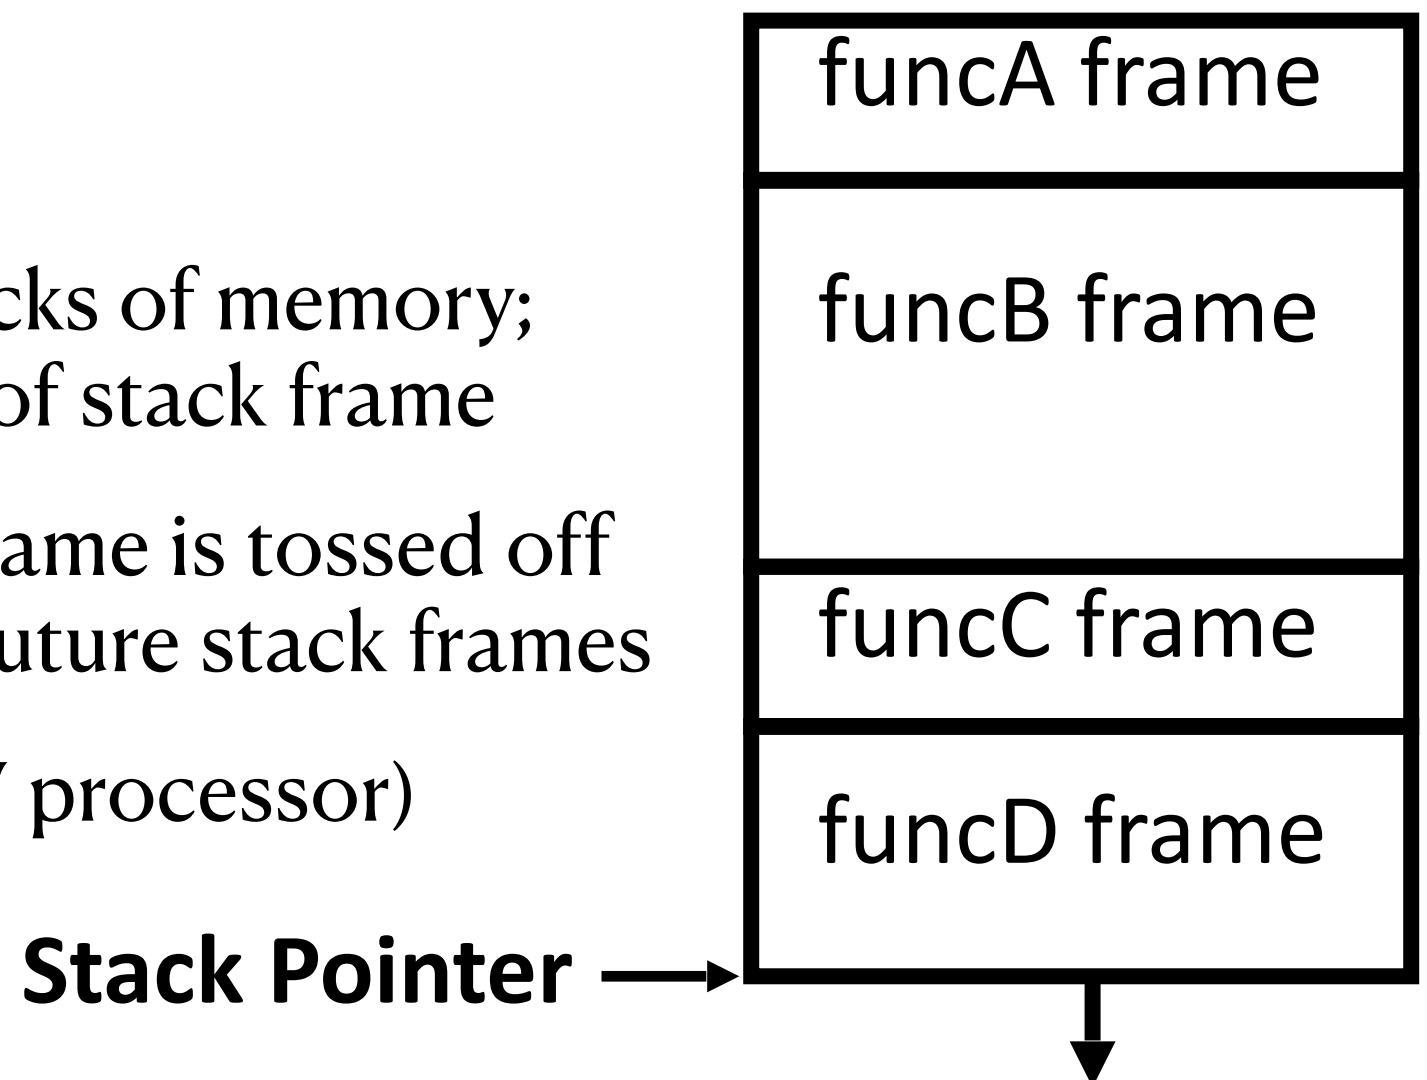
- If declared outside a function, allocated in “static” storage
- If declared inside function, allocated on the “stack” and freed when function returns
  - `main()` is treated like a function
- For the above two types, the memory management is automatic
  - Don't need to deallocate when no longer using them
  - A variable does not exist anymore once a function ends!

```
int myGlobal;  
main() {  
    int myTemp;  
}
```

# The Stack

- Every time a function is called, a new “stack frame” is allocated on the stack
- Stack frame includes:
  - Return address (who called me?)
  - Arguments
  - Space for local variables
- Stack frames contiguous blocks of memory; stack pointer indicates start of stack frame
- When function ends, stack frame is tossed off the stack; frees memory for future stack frames
- Details covered later (RISC-V processor)

```
funcA() { funcB(); }  
funcB() { funcC(); }  
funcC() { funcD(); }
```



# Stack Animation

- Last In, First Out (LIFO) data structure

```
main ()
{ a(0);
}

void a (int m)
{ b(1);
}


void b (int n)
{ c(2);
}

void c (int o)
{ d(3);
}

void d (int p)
{
}
```

**Stack Pointer**

Stack  
grows  
down




```
_main:
sub sp, sp, #48
... ..
mov w8, #1234
stur w8, [x29, #-8]
mov w8, #4321
stur w8, [x29, #-12]
ldur w8, [x29, #-8]
ldur w9, [x29, #-12]
add w8, w8, w9
... ..
add sp, sp, #48
ret
```

# Passing Pointers into the Stack

- It is fine to pass a pointer to stack space further down.

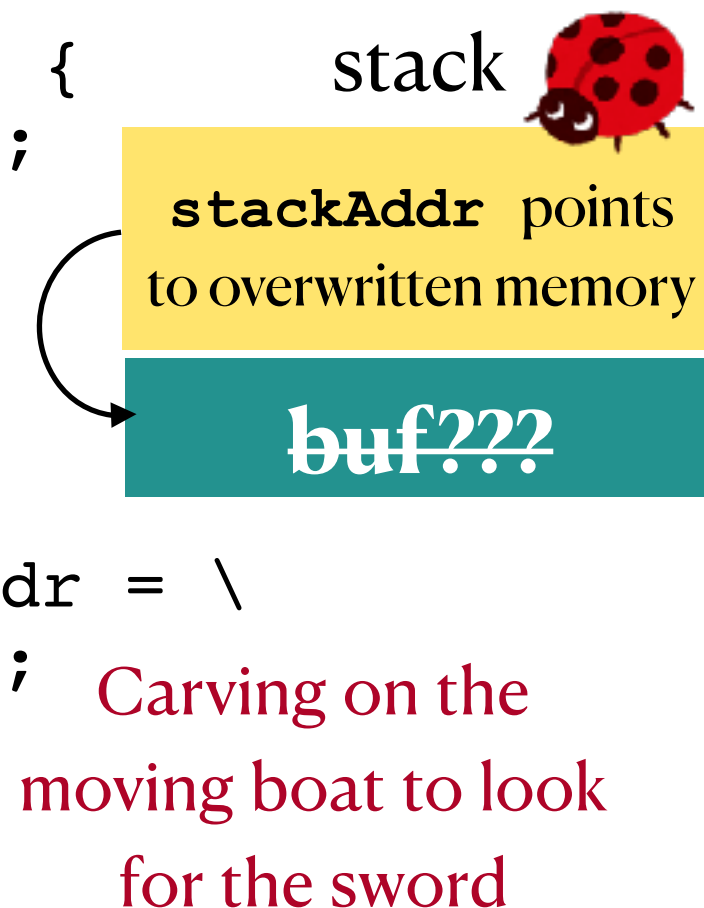
```
#define BUFSIZE 256
int main() {
    ...
    char buf[BUFSIZE];
    load_buf(buf, BUFSIZE);
    ...
}
```



- However, it is bad to return a pointer to something in the **stack**!
- Memory will be overwritten when other functions called!
- So your data would no longer exist, and writes can overwrite key pointers, causing crashes!

```
char *make_buf() {
    char buf[50];
    return buf;
}
```

```
int main() {
    ...
    char *stackAddr = \
        make_buf();
    foo();
    ...
}
```



Solve with slides to come ...

# Managing the Heap

- The heap is **dynamic** memory – memory that can be allocated, resized, and freed during program runtime.
  - Useful for persistent memory across function calls
  - But biggest source of pointer bugs, memory leaks, ...
- Large pool of memory, not allocated in contiguous order
  - Back-to-back requests for heap memory could result in blocks very far apart
- C supports four functions for heap management:
  - **malloc()**      allocate a block of uninitialized memory
  - **calloc()**      allocate a block of zeroed memory
  - **free()**          free previously allocated block of memory
  - **realloc()**      change size of previously allocated block (might move)
- Read-more: <http://web.archive.org/web/20030222051144/http://home.earthlink.net/~bobbitts/c89.txt> section 4.10.3 memory management functions

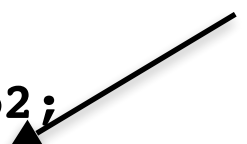
# Managing the Heap

- **void \*malloc(size\_t n):**
  - Allocate a block of uninitialized memory
  - **n** is an integer, indicating size of allocated memory block in bytes
  - **size\_t** is an unsigned integer type big enough to “count” memory bytes
  - **sizeof** returns size of given type in bytes, produces more portable code
  - Returns **void\*** pointer to block; **NULL** return indicates no more memory; **always check** for return **NULL** (**if (ip)**)
  - Think of pointer as a handle that describes the allocated block of memory; Additional control information stored in the heap around the allocated block! (Including size, etc.)

*“Cast” operation, changes type of a variable.*

- Examples:

```
int *ip1, *ip2;  
ip1 = (int *) malloc(sizeof(int));  
ip2 = (int *) malloc(20*sizeof(int)); //allocate an array of 20 ints.
```



```
typedef struct { ... } TreeNode;  
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));
```

Assuming size of objects can lead to misleading, unportable code. Use **sizeof()**!

# Managing the Heap

- **void free(void \*p) :**
  - Releases memory allocated by **malloc()**
  - **p** is pointer containing the address originally returned by **malloc()**

```
int *ip;
ip = (int *) malloc(sizeof(int));
... ..
free((void*) ip); /* Can you free(ip) after ip++ ? */

typedef struct {...} TreeNode;
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));
... ..
free((void *) tp);
```
  - When you free memory, you must be sure that you pass the original address returned from **malloc()** to **free()**; Otherwise, system exception (or worse)!

# Managing the Heap

- **`void *realloc(void *p, size_t size):`**
    - Returns new address of the memory block.
      - In doing so, it may need to copy all data to a new location.
- `realloc(NULL, size); // behaves like malloc`**  
**`realloc(ptr, 0); // behaves like free, deallocates heap block`**
- **Always check** for return NULL

```
int *ip; ip = (int *) malloc(10*sizeof(int));  
... .. /* check for NULL */  
ip = (int *) realloc(ip, 20*sizeof(int));  
/* contents of first 10 elements retained */  
... .. /* check for NULL */  
realloc(ip, 0); /* equivalent to free(ip); */
```

Keep track of  
this, since it  
might change.



# Summary

- Code, static storage are easy: they never grow or shrink
- Stack space is relatively easy: stack frames are created and destroyed in last-in, first-out (LIFO) order, avoid “dangling references”
- Managing the heap is tricky:
  - Memory can be allocated/deallocated at any time
  - “Memory leak”: If you forget to deallocate memory
  - “Use after free”: If you use data after calling free
  - “Double free”: If you call free 2x on same memory



# Using Dynamic Memory—Linked List

```
typedef struct Node
{
    int val;
    struct Node *next;
} node;
```

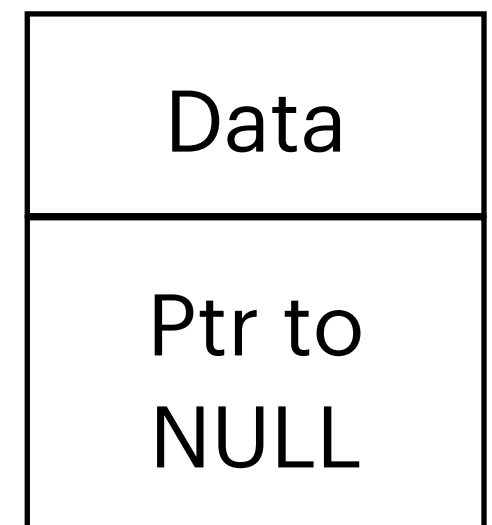

```
node * head = NULL;
head = (node *) malloc(sizeof(node));
if(head == NULL){
    return 1;
}
head -> val = 1;
head -> next = NULL;
```

Create the first node

The first node



The last node



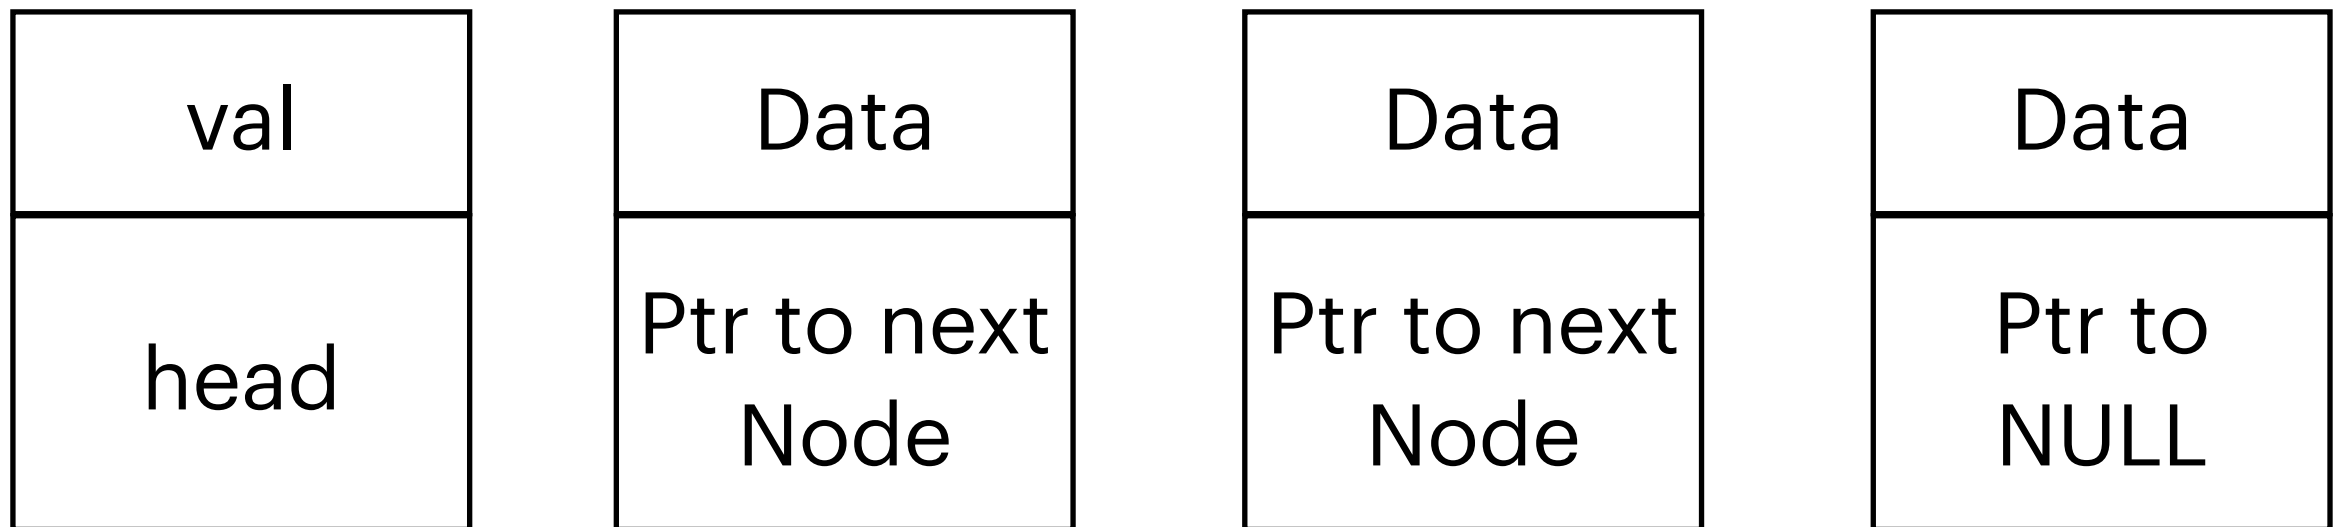
↑  
Ptr to head

# Using Dynamic Memory—Iterate

```
typedef struct Node
{
    int val;
    struct Node *next;
} node;
```

```
void print_list(node *head){
    node * current = head;
    while (current != NULL){
        printf("%d\t", current -> val);
        current = current -> next;
    }
    printf("\n");
}
```

The first node




↑  
Ptr to current node

# Using Dynamic Memory—Push

```
typedef struct Node
{
    int val;
    struct Node *next;
} node;
```

```
void push_node(node ** head, int val){
    node * new_node;
    new_node = (node *) malloc (sizeof
(node));
    new_node -> val = val;
    new_node -> next = *head;
    *head = new_node;
    printf("Node %d push succeeds!\n",
(*head) -> val);
}
```



The first node



# Using Dynamic Memory—Remove

```
typedef struct Node
{
    int val;
    struct Node *next;
} node;
```

The first node



Ptr to cur. node

```
int remove_last(node * head) {
    int retval = 0;

    if (head->next == NULL) {
        retval = head->val;
        free(head);
        return retval;
    }



    node * current = head;
    while (current->next->next != NULL) {
        current = current->next;
    }

    retval = current->next->val;
    free(current->next);
    current->next = NULL;
    printf("%d is removed.\n", retval);
    return retval;
}
```

# How are Malloc/Free implemented?

- Underlying operating system allows **malloc** library to ask for large blocks of memory to use in heap (e.g., using Unix **sbrk()** call)
- C standard **malloc** library creates data structure inside unused portions to track free space


# Simple Slow Malloc Implementation



Malloc library creates linked list of empty blocks (one block initially)



## First allocation chews up space from start of free space




After many mallocs and frees, have potentially long linked list of odd-sized blocks  
Frees link block back onto linked list – might merge with neighboring free space

# Faster malloc implementations

- Keep separate pools of blocks for different sized objects
- “Buddy allocators” always round up to power-of-2 sized chunks to simplify finding correct size and merging neighboring blocks:



# Power-of-2 “Buddy Allocator”



# Malloc Implementations

- All provide the same library interface, but can have radically different implementations
- Uses headers at start of allocated blocks and space in unallocated memory to hold **malloc**'s internal data structures
- Rely on programmer remembering to free with same pointer returned by **malloc**
- Rely on programmer not messing with internal data structures accidentally!

# Agenda

- C Memory Management
- C Bugs

# Common Memory Problems

- Using uninitialized values
- Using memory that you don't own
  - Deallocated stack or heap variable
  - Out-of-bounds reference to stack or heap array
  - Using NULL or garbage data as a pointer
- Improper use of free/realloc by messing with the pointer handle returned by malloc/calloc
- Memory leaks (you allocated something you forgot to later free)

# Using Memory You Don't Own

- What is wrong with this code?
- Using pointers beyond the range that had been malloc'd
  - May look obvious, but what if mem refs had been result of pointer arithmetic that erroneously took them out of the allocated range?

```
int *ipr, *ipw;
void ReadMem() {
    int i, j;
    ipr = (*int) malloc(4 * sizeof(int));
    i = *(ipr - 1000); j = *(ipr + 1000);
    free(ipr);
}
```

```
void WriteMem() {
    ipw = (*int) malloc(5 * sizeof(int));
    *(ipw - 1000) = 0; *(ipw + 1000) = 0;
    free(ipw);
}
```

# Faulty Heap Management

- What is wrong with this code?

```
int *pi;  
void foo() {  
    pi = malloc(8*sizeof(int));  
    ...  
    free(pi);  
}
```

```
void main() {  
    pi = malloc(4*sizeof(int));  
    foo();  
    ...  
}
```

# Faulty Heap Management

- Memory leak: more mallocs than frees

```
int *pi;
void foo() {
    pi = malloc(8*sizeof(int));
    /* Allocate memory for pi */
    /* Oops, leaked the old memory pointed to by pi */
    ...
    free(pi); /* foo() is done with pi, so free it */
}

void main() {
    pi = malloc(4*sizeof(int));
    foo(); /* Memory leak: foo leaks it */
    ...
}
```

# Faulty Heap Management

- What is wrong with this code?

```
int *plk = NULL;
void genPLK() {
    plk = malloc(2 * sizeof(int));
    ... ..
    plk++;
}
```



# Faulty Heap Management

- Potential memory leak – handle has been changed, do you still have copy of it that can correctly be used in a later free?

```
int *plk = NULL;
void genPLK() {
    plk = malloc(2 * sizeof(int));
    ... ..
    plk++;
}
```

# Faulty Heap Management

- What is wrong with this code?

```
void FreeMemX() {  
    int fnh = 0;  
    free(&fnh);  
}
```

```
void FreeMemY() {  
    int *fum = malloc(4 * sizeof(int));  
    free(fum+1);  
    free(fum);  
    free(fum);  
}
```

# Faulty Heap Management

- Can't free non-heap memory; Can't free memory that hasn't been allocated

```
void FreeMemX() {  
    int fnh = 0;  
    free(&fnh);  
}
```

```
void FreeMemY() {  
    int *fum = malloc(4 * sizeof(int));  
    free(fum+1);  
    free(fum);  
    free(fum);  
}
```

# Using Memory You Haven't Allocated

- What is wrong with this code?

```
void StringManipulate() {  
    const char *name = "Safety Critical";  
    char *str = malloc(10);  
    strncpy(str, name, 10);  
    str[10] = '\0';  
    printf("%s\n", str);  
}
```

# Using Memory You Haven't Allocated

- Reference beyond array bounds

```
void StringManipulate() {  
    const char *name = "Safety Critical";  
    char *str = malloc(10);  
    strncpy(str, name, 10);  
    str[10] = '\0';  
    /* Write Beyond Array Bounds */  
    printf("%s\n", str);  
}
```

# Using Memory You Don't Own

- What's wrong with this code?

```
char *append(const char* s1, const char *s2) {  
    const int MAXSIZE = 128;  
    char result[128];  
    int i=0, j=0;  
    for (j=0; i<MAXSIZE-1 && j<strlen(s1); i++,j++) {  
        result[i] = s1[j];  
    }  
    for (j=0; i<MAXSIZE-1 && j<strlen(s2); i++,j++) {  
        result[i] = s2[j];  
    }  
    result[++i] = '\0';  
    return result;  
}
```

# Using Memory You Don't Own

- Beyond stack read/write

```
char *append(const char* s1, const char *s2) {  
    const int MAXSIZE = 128;  
    char result[128];  
    int i=0, j=0;  
    for (j=0; i<MAXSIZE-1 && j<strlen(s1); i++,j++) {  
        result[i] = s1[j];  
    }  
    for (j=0; i<MAXSIZE-1 && j<strlen(s2); i++,j++) {  
        result[i] = s2[j];  
    }  
    result[++i] = '\0';  
    return result;  
}
```

result is a local array name –  
stack memory allocated

Function returns pointer to stack  
memory – won't be valid after function  
returns

# Using Memory You Don't Own

- What is wrong with this code?

```
typedef struct node {  
    struct node* next;  
    int val;  
} Node;
```

```
int findLastNodeValue(Node* head) {  
    while (head->next != NULL) {  
        head = head->next;  
    }  
    return head->val;  
}
```



# Using Memory You Don't Own

- Following a NULL pointer to mem addr 0!

```
typedef struct node {  
    struct node* next;  
    int val;  
} Node;
```

```
int findLastNodeValue(Node* head) {  
    while (head->next != NULL) {  
        head = head->next;  
    }  
    return head->val;  
}
```

# Using Memory You Don't Own

- What is wrong with this code?

```
int* init_array(int *ptr, int new_size) {  
    ptr = realloc(ptr, new_size*sizeof(int));  
    memset(ptr, 0, new_size*sizeof(int));  
    return ptr;  
}
```

```
int* fill_fibonacci(int *fib, int size) {  
    int i;  
    init_array(fib, size);  
    /* fib[0] = 0; */ fib[1] = 1;  
    for (i=2; i<size; i++)  
        fib[i] = fib[i-1] + fib[i-2];  
    return fib;  
}
```

# Using Memory You Don't Own

- Improper matched usage of mem handles

```
int* init_array(int *ptr, int new_size) {  
    ptr = realloc(ptr, new_size*sizeof(int));  
    memset(ptr, 0, new_size*sizeof(int));  
    return ptr;  
}
```

Remember: **realloc** may move entire block

```
int* fill_fibonacci(int *fib, int size) {  
    int i;  
    /* oops, forgot: fib = */ init_array(fib, size);  
    /* fib[0] = 0; */ fib[1] = 1;  
    for (i=2; i<size; i++)  
        fib[i] = fib[i-1] + fib[i-2];  
    return fib;  
}
```

What if array is moved to new location?

# Summary

- All data/program is in memory
  - Each memory location has an address to use to refer to it and a value stored in it
- Pointer is a C version (abstraction) of a data address
  - \* “follows” a pointer to its value
  - & gets the address of a value
  - Arrays and strings are implemented as variations on pointers
- C is an efficient language, but leaves safety to the programmer
  - Variables not automatically initialized
  - Use pointers with care: they are a common source of bugs in programs

# Summary

- C has three main memory segments in which to allocate data:
  - Static Data: Variables outside functions
  - Stack: Variables local to function
  - Heap: Objects explicitly malloc-ed/free-d.
- Heap data is biggest source of bugs in C code