



信息科学与技术学院

School of Information Science and Technology

CS 110

Computer Architecture

Intro to RISC-V I

Instructors:

Siting Liu & Chundong Wang

Course website: <https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html>

School of Information Science and Technology (SIST)

ShanghaiTech University

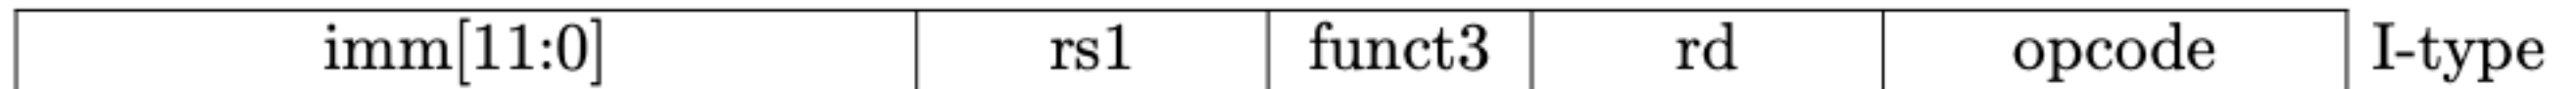
2023/2/6

Course Info

- Lab 3 will be released after class (10 a.m.), get yourself prepared before going to lab sessions!
- Our project 1.1 will be available this weekend, and will be marked in lab sessions. Deadline March 13th.
- Next week discussion on RISC-V related materials and assembly coding.

Assembly Instructions

- Different types of instructions



- I-type
 - Register-Immediate type
 - Has two operands (one accessed from source register, another a constant/immediate) and one output (saved to destination register)
 - Can do arithmetic/logic/load from main memory/**jump (covered later)**

RV32I I-type Arithmetic

- Syntax of instructions: assembly language

- Addition: `addi rd, rs1, imm`

Adds `imm` to `rs1`, stores the result to `rd`, and `imm` is a signed number.

- Example: `addi x5, x4, 10`
`addi x6, x4, -10`

- Similarly, `andi/ori/xori/slti/sltiu`

- All the `imm`'s are sign-extended

- `slli/srli/srai` are special (defined in RV64I), and can be extended to RV32I usage (RTFM)

Registers

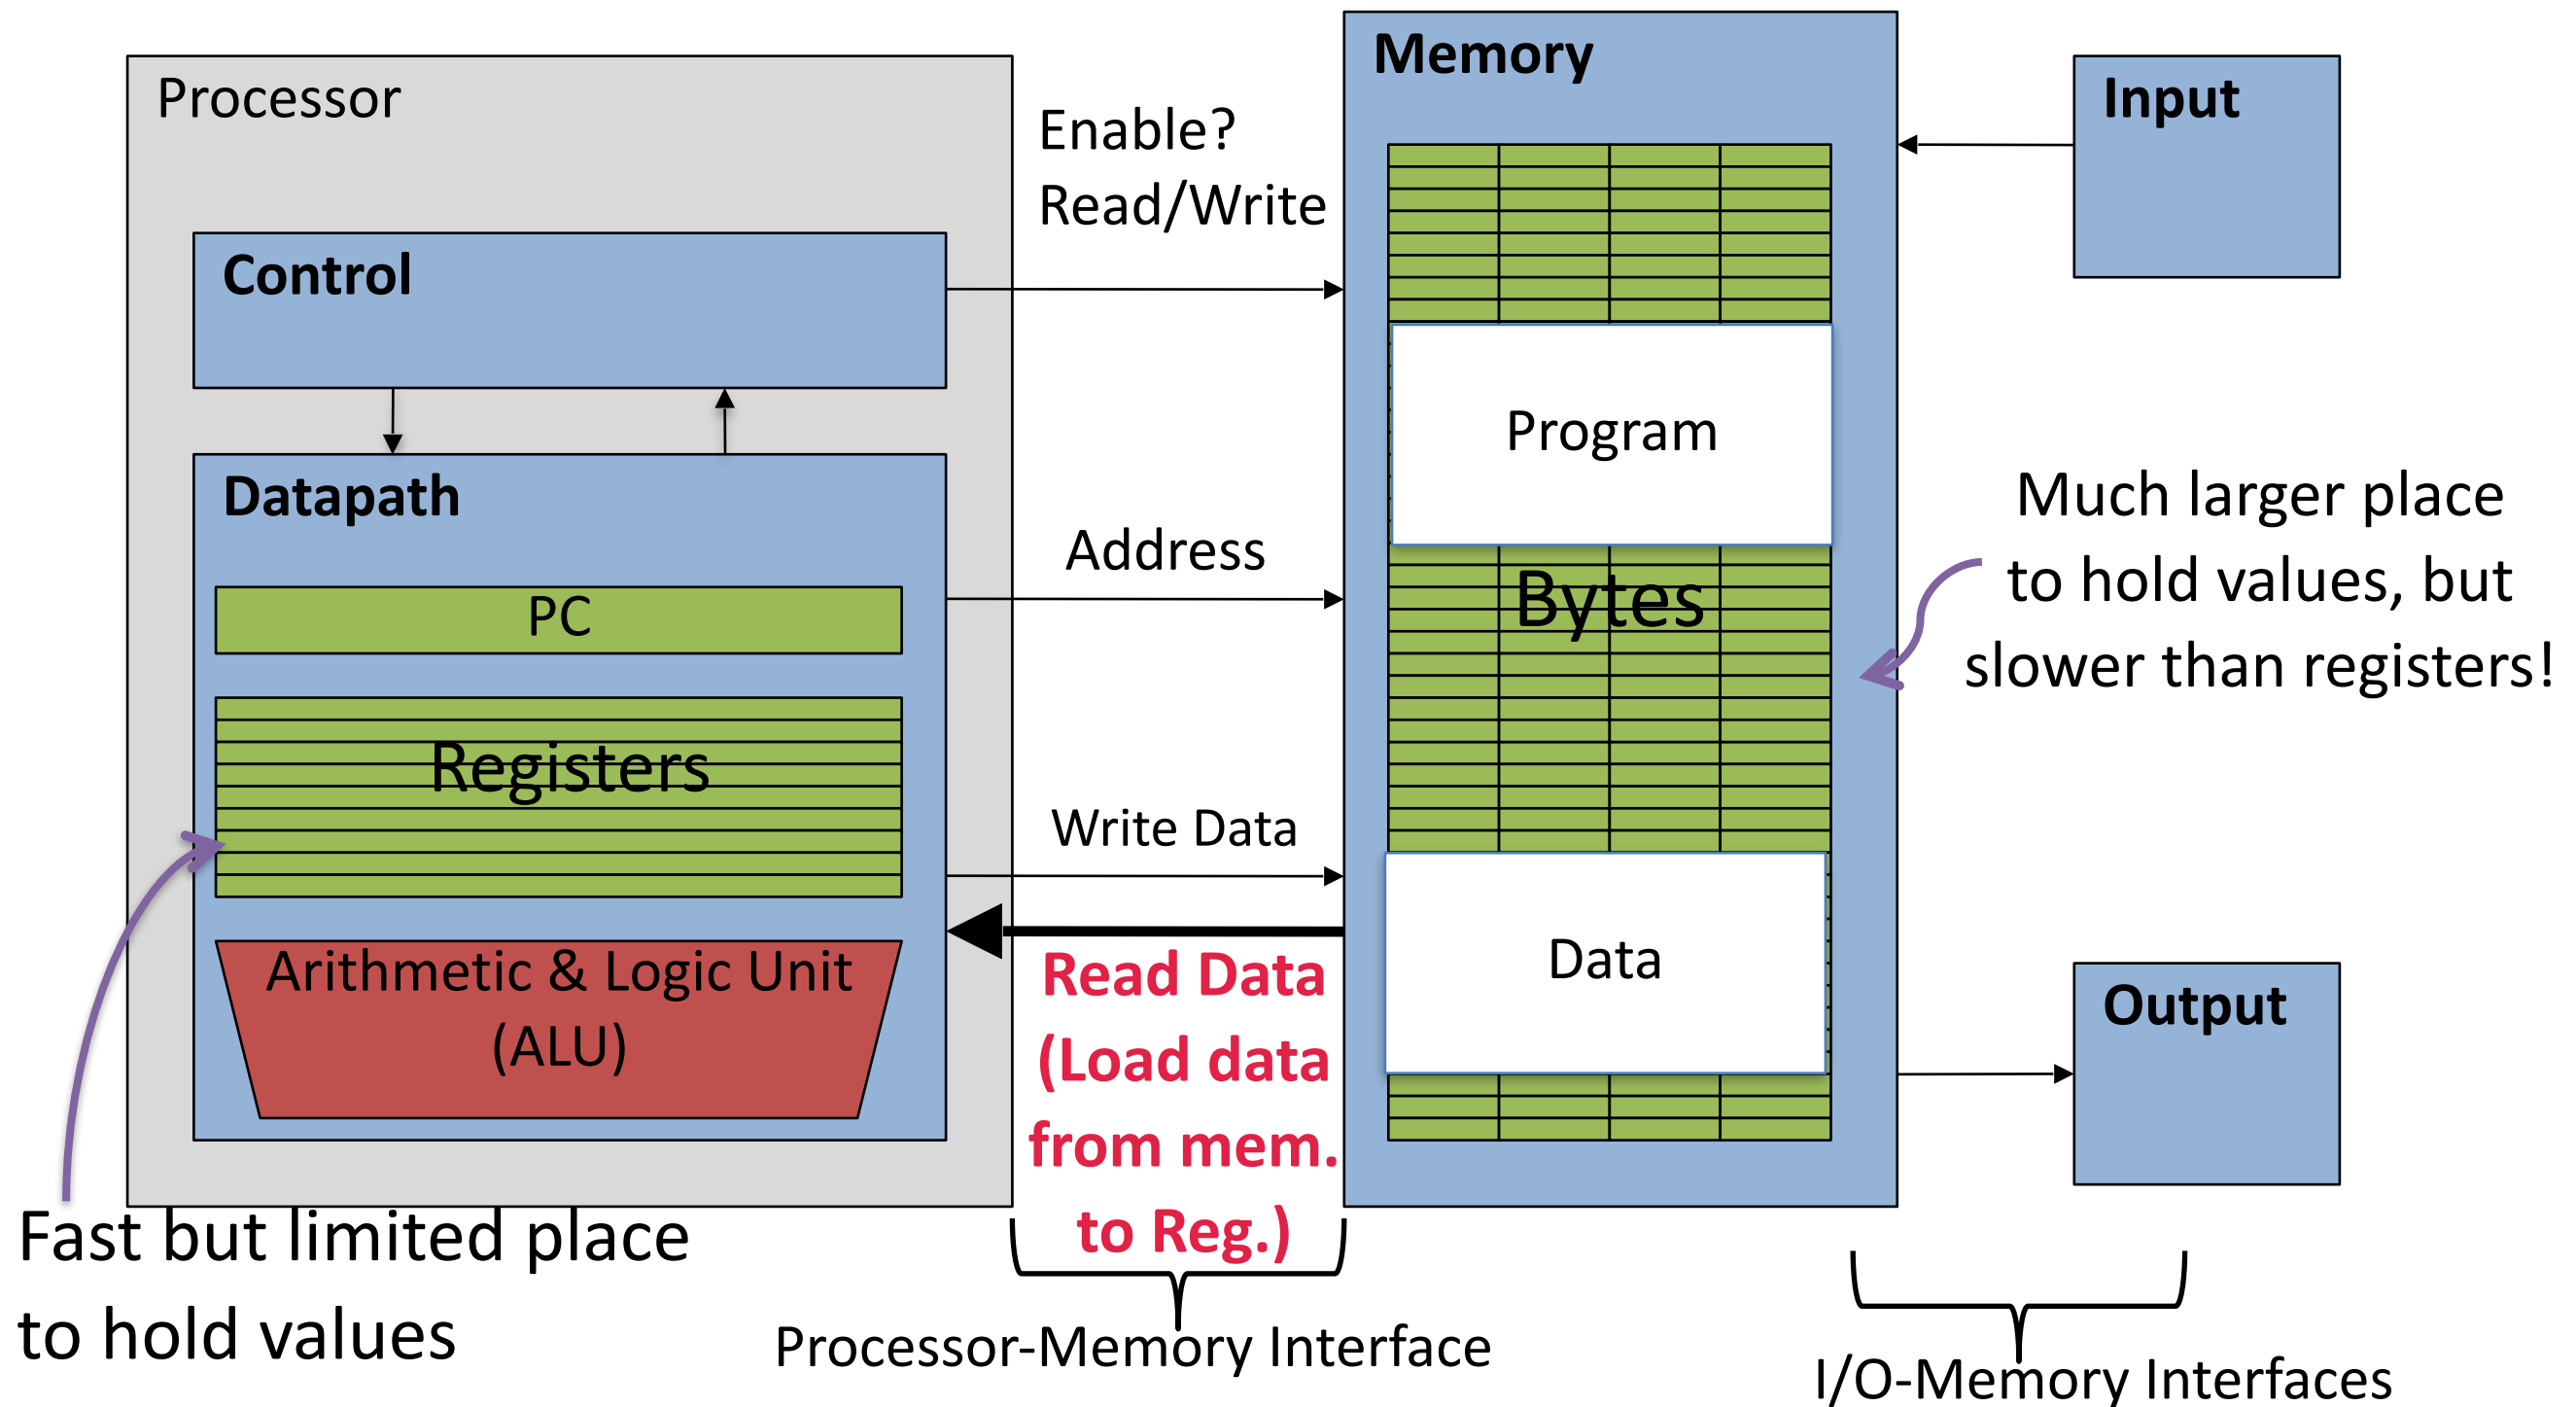
0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x3	x4
	x5
	x6
	x7

RV32I Arithmetic/Logic Test

- `addi x1, x2, -1`
 - `or x2, x2, x1`
 - `add x3, x1, x2`
 - `slt x4, x3, x1`
 - `sra x5, x3, x4`
 - `sub x0, x5, x4`
- Register zero (**x0**) is 'hard-wired' to 0;
 - By **convention** RISC-V has a specific **no-op** instruction...
 - **`add x0 x0 x0`**
 - You may need to replace code later: No-ops can fill space, align data, and perform other options
 - Practical use in jump-and-link operations (covered later)

Registers	
0	x0/zero
0	x1
0	x2
0	x3
0	x4
0	x5
0	x6
0	x7

RV32I I-type Load



Big Endian vs. Little Endian

Big-endian and little-endian from Jonathan Swift's *Gulliver's Travels*

- The order in which BYTES are stored in memory
- Bits always stored as usual. (E.g., 0xC2=0b 1100 0010)

Consider the number 1025 as we normally write it:

BYTE3 BYTE2 BYTE1 BYTE0
00000000 00000000 0000100 00000001

Big Endian

ADDR3 ADDR2 ADDR1 ADDR0
BYTE0 BYTE1 BYTE2 BYTE3
00000001 0000100 00000000 00000000

Examples

Names in the West (e.g. Siting, Liu)

Java Packages: (e.g. org.mypackage.HelloWorld)

Dates done correctly ISO 8601 YYYY-MM-DD
(e.g. 2020-03-22)

Eating Pizza crust first

Unix file structure (e.g., /usr/local/bin/python)

"Network Byte Order": most network protocols

IBM z/Architecture; very old Macs

Little Endian

ADDR3 ADDR2 ADDR1 ADDR0
BYTE3 BYTE2 BYTE1 BYTE0
00000000 00000000 0000100 00000001

Examples

Names in China (e.g. LIU Siting)

Internet names (e.g. sist.shanghaitech.edu.cn)

Dates written in England DD/MM/YYYY
(e.g. 22/03/2020)

Eating Pizza skinny part first (the normal way)

CANopen

Intel x86; **RISC-V** (can also support big-endian)

Assembly Instructions—Load

- RV32I is a load-store architecture, where only load and store instructions access memory and arithmetic instructions only operate on CPU registers.

imm[11:0]	rs1	funct3	rd	opcode	I-type
-----------	-----	--------	----	--------	--------

- lw** *rd*, **imm**(*rs1*) : Load word at *addr.* to register *rd*
 $\text{addr.} = (\text{number in rs1}) + \text{imm}$

- Example

`lw x1, 12(x4)`

$\text{addr.} = 4 + 12 = (10)_{\text{HEX}}$

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x4	x4
	x5
	x6
	x7

Registers

56	34	23	01	3c
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	...
34	12	cd	ab	c
56	34	23	01	8
34	12	cd	ab	4
				0

Main memory

Assembly Instructions—Load

- RV32I is a load-store architecture, where only load and store instructions access memory and arithmetic instructions only operate on CPU registers.

imm[11:0]	rs1	funct3	rd	opcode	I-type
-----------	-----	--------	----	--------	--------

- `lw rd, imm(rs1)` : Load word at addr. to register rd

addr. = (number in rs1) + imm

- Example

```
lw x1, 12(x4)
```

addr. = 4 + 12 = (10)_{HEX}

- C code example

```
int A[100];
```

```
/*assume &A[0] = 4*/
```

```
G = A[3];
```

```
/*load G to x1/
```

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x4	x4
	x5
	x6
	x7

Registers

56	34	23	01	3c
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	...
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	10
56	34	23	01	C
34	12	cd	ab	8
56	34	23	01	4
34	12	cd	ab	0

Main memory

Assembly Instructions—Load

- RV32I is a load-store architecture, where only load and store instructions access memory and arithmetic instructions only operate on CPU registers.

imm[11:0]	rs1	funct3	rd	opcode	I-type
-----------	-----	--------	----	--------	--------

- `lb/lbu rd, imm(rs1)` : Load signed/unsigned byte at `addr.` to register `rd`

$\text{addr.} = (\text{number in rs1}) + \text{imm}$

- Example

`lb x1, 12(x4)`

$\text{addr.} = 4 + 12 = (10)_{\text{HEX}}$

`lbu x1, 12(x4)`

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x4	x4
	x5
	x6
	x7

Registers

56	34	23	01	3c
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	...
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	10
56	34	23	01	c
34	12	cd	ab	8
56	34	23	01	4
34	12	cd	ab	0

Main memory

Assembly Instructions—Load

- RV32I is a load-store architecture, where only load and store instructions access memory and arithmetic instructions only operate on CPU registers.

imm[11:0]	rs1	funct3	rd	opcode	I-type
-----------	-----	--------	----	--------	--------

- `lh/lhu rd, imm(rs1)` : Load signed/unsigned halfword at `addr.` to register `rd` (similar to `lb/lbu`)

$\text{addr.} = (\text{number in rs1}) + \text{imm}$

- Example

`lh x1, 12(x4)`

$\text{addr.} = 4 + 12 = (10)_{\text{HEX}}$

`lhu x1, 12(x4)`

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x4	x4
	x5
	x6
	x7

Registers

56	34	23	01	3c
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	...
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	10
34	12	cd	ab	C
56	34	23	01	8
34	12	cd	ab	4
56	34	23	01	0
34	12	cd	ab	0

Main memory

Assembly Instructions—S-Type Store

- RV32I is a load-store architecture, where only load and store instructions access memory and arithmetic instructions only operate on CPU registers.

imm[11:5]	rs2	rs1	funct3	imm[4:0]	opcode	S-type
-----------	-----	-----	--------	----------	--------	--------

- `sw rs2, imm(rs1)` : Store word at rs2 to memory addr.

$\text{addr.} = (\text{number in rs1}) + \text{imm}$

- Example

```
sw x1, 12(x4)
```

$\text{addr.} = 4 + 12 = (10)_{\text{HEX}}$

- C code example

```
int A[100];
```

```
/* &A[0] => x4 */
```

```
A[3] = h;
```

```
/* h in rs2 => x1 */
```

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x4	x4
	x5
	x6
	x7

Registers

56	34	23	01	3c
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	
56	34	23	01	...
34	12	cd	ab	
56	34	23	01	
34	12	cd	ab	10
56	34	23	01	c
34	12	cd	ab	8
56	34	23	01	4
34	12	cd	ab	0

Main memory

Assembly Instructions—S-Type Store

- RV32I is a load-store architecture, where only load and store instructions access memory and arithmetic instructions only operate on CPU registers.

imm[11:5]	rs2	rs1	funct3	imm[4:0]	opcode	S-type
-----------	-----	-----	--------	----------	--------	--------

- `sw rs2, imm(rs1)` : Store word at rs2 to memory addr .

$\text{addr.} = (\text{number in rs1}) + \text{imm}$

- Example

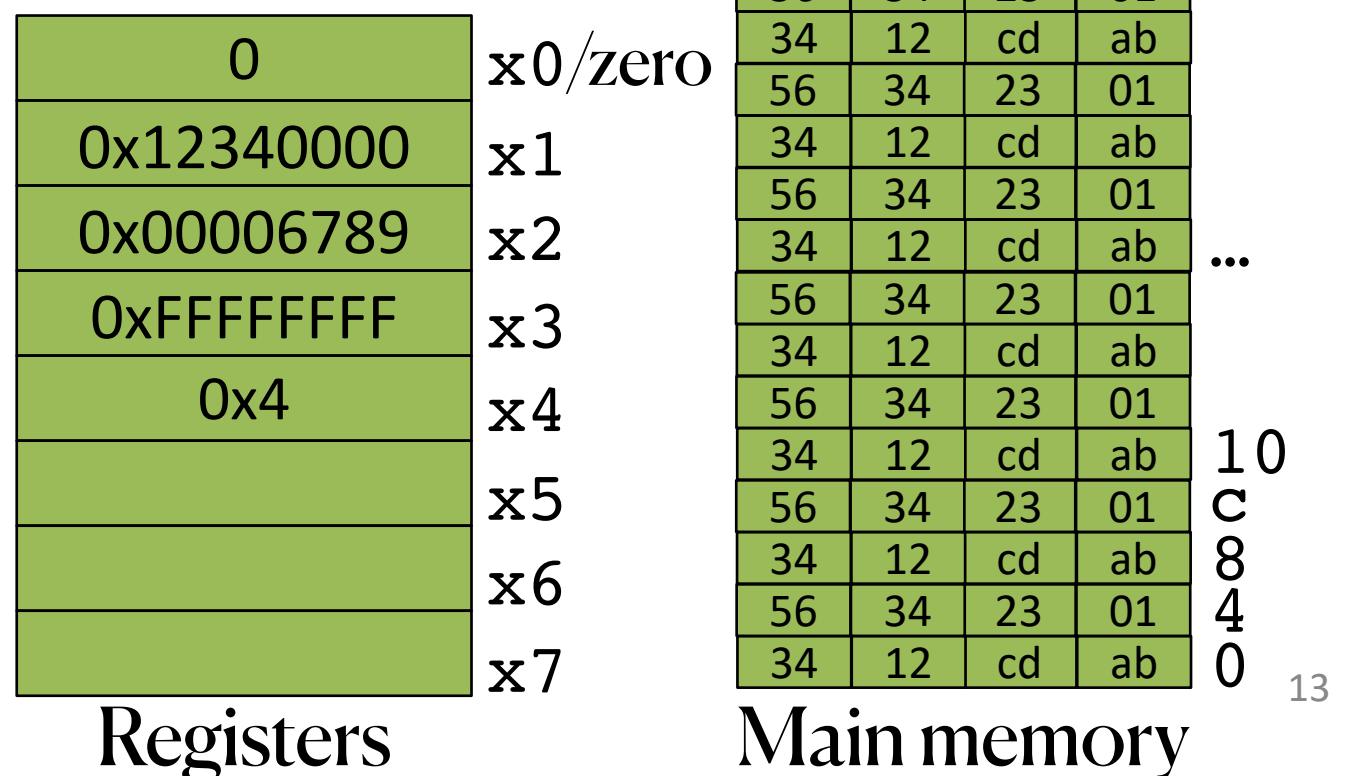
`sw x1, 12(x4)`

$\text{addr.} = 4 + 12 = (10)_{\text{HEX}}$

- Similarly,

`sh` : Store lower 16 bits at rs2

`sb` : Store lower 8 bits at rs2



Memory Alignment

- RISC-V *does not require* that integers be word aligned...
 - But it can be very very bad if you don't make sure they are...
- Consequences of unaligned integers
 - Slowdown: The processor is allowed to be a lot slower when it happens
 - In fact, a RISC-V processor may natively only support aligned accesses, and do unaligned-access in software!
An unaligned load could take hundreds of times longer!
 - Lack of atomicity: The whole thing doesn't happen at once... can introduce lots of very subtle bugs
- So in practice, RISC-V *recommends* integers to be aligned on 4- byte boundaries; halfword 2-byte boundaries

Question! What's in x12?

addi x11,x0,0x4F6

sw x11,0(x5)

lb x12,1(x5)

A :

0x0

B :

0x4

C :

0x6

D :

0xF

E :

0xFFFFFFFF

Question! What's in x12?

addi x11,x0,0x85F6

sw x11,0(x5)

lb x12,1(x5)

A:	0x8
B:	0x85
C:	0xC
D:	0xBC
E:	0xFFFFFFFF85
F:	0xFFFFFFFFF8
G:	0xFFFFFFFFFC
H:	0xFFFFFFFFBC

Summary

- RISC-V ISA basics: (32 registers, referred to as $x0-x31$, $x0=0$)
- Simple is better
- One instruction (simple operation) per line (RISC-V assembly)
- Fixed-length instructions (for RV32I)
- 6 types of instructions (depending on their format)
- Instructions for arithmetics, logic operations, register-memory data exchange (load/store word/halfword/byte)
- RISC-V is little-endian
- Load-store architecture



信息科学与技术学院

School of Information Science and Technology

CS 110

Computer Architecture

Intro to RISC-V II

Computer Decision Making

Instructors:

Siting Liu & Chundong Wang

Course website: [https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/
Spring-2023/index.html](https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html)

School of Information Science and Technology (SIST)

ShanghaiTech University

2023/2/6

Computer Decision Making—Branch

- Normal operation: execute instructions in sequence
- In C: `if/while/for`-statement; function call
- RISV-V provides conditional branch (B-type) & unconditional jump (j)

imm[12]	imm[10:5]	rs2	rs1	funct3	imm[4:1]	imm[11]	opcode	B-type
---------	-----------	-----	-----	--------	----------	---------	--------	--------

- RISC-V: similar `if`-statement instruction

`beq rs1, rs2, L(imm/label)`

means: go to statement labeled if (value in `rs1`) == (value in `rs2`);

otherwise, go to next statement

- `beq` stands for branch if equal
- Similarly, `bne` for branch if not equal

Computer Decision Making—Branch

- Example:

`beq rs1, rs2, L(imm/label)`

- C code

```
int main(void) {  
    int i=5;  
    if (i!=6){  
        i++;  
    }  
    else i--;  
    return 0;  
}
```

- Assembly

```
addi x2, x0, 5  
addi x3, x0, 6  
bne x2, x3, L1  
beq x2, x3, L2  
L1:addi x2, x2, 1  
    ret (kind of jump)  
L2:addi x2, x2, -1  
    ret
```

- Label can also point to data (more in discussion)

Computer Decision Making—Branch

- Example:
`beq rs1, rs2, L(imm/label)`

- C code

```
int main(void) {  
    int i=5;  
    if (i!=6){  
        i++;  
    }  
    else i--;  
    return 0;  
}
```

- Assembly (real stuff in ARM64)

```
    mov w8, #5  
Ltmp3:  
    .loc1 10 9 is_stmt 0  
    subs w8, w8, #6  
    b.eq LBB0_2  
    b     LBB0_1  
LBB0_1:  
Ltmp4:  
    .loc1 11 10 is_stmt 1  
    ldr w8, [sp, #8]  
    add w8, w8, #1  
    str w8, [sp, #8]  
    .loc1 12 5  
    b     LBB0_3  
Ltmp5:  
LBB0_2:  
    .loc1 13 11  
    ldr w8, [sp, #8]  
    subs w8, w8, #1  
    str w8, [sp, #8]  
    b     LBB0_3  
Ltmp6:  
LBB0_3:  
    .loc1 0 11 is_stmt 0  
    mov w0, #0  
    .loc1 14 5 is_stmt 1  
    add sp, sp, #16  
    ret
```

Computer Decision Making—Branch

- Normal operation: execute instructions in sequence
- In programming languages: `if/while/for`-statement
- RISV-V provides conditional branch & unconditional jump

imm[12]	imm[10:5]	rs2	rs1	funct3	imm[4:1]	imm[11]	opcode	B-type
---------	-----------	-----	-----	--------	----------	---------	--------	--------

- RISC-V: `if`-statement instructions are

`blt/bltu/bge/bgeu rs1, rs2, L(imm/label)`

means: go to statement labeled L1 if (value in `rs1`) $</\geq$ (value in `rs2`)
using signed/unsigned comparison; otherwise, go to next statement

C Loop Mapped to RISC-V Assembly

```
int A[20];  
int sum = 0;  
for (int i=0; i < 20; i++)  
    sum += A[i];
```

```
# Assume x8 holds pointer to A  
# Assign x10=sum  
add x9, x8, x0 # x9=&A[0]  
add x10, x0, x0 # sum=0  
add x11, x0, x0 # i=0  
addi x13, x0, 20 # x13=20  
Loop:  
    bge x11, x13, Done  
    lw x12, 0(x9) # x12=A[i]  
    add x10, x10, x12 # sum+=  
    addi x9, x9, 4 # &A[i+1]  
    addi x11, x11, 1 # i++  
    j Loop  
Done:
```

Optimization

- The simple translation is sub-optimal!
 - Inner loop is now 4 instructions rather than 7
 - And only 1 branch/jump rather than two: Because first time through is always true so can move check to the end!
- The compiler will often do this automatically for optimization
 - See that i is only used as an index in a loop

```
# Assume x8 holds pointer to A
# Assign x10=sum
add  x10, x0, x0 # sum=0
add  x11, x8, x0 # ptr = A
addi x12, x11, 80 # end = A + 80
Loop:
    lw   x13, 0(x11) # x13 = *ptr
    add  x10, x10, x13 # sum += x13
    addi x11, x11, 4 # ptr++
    blt  x11, x12, Loop: # ptr < end
```

- This optimization is not required
- Line by line translation is good
- Correctness first, performance second

Arrays and Pointers

```
int i;  
int array[10];  
  
for (i = 0; i < 10; i++)  
{  
    array[i] = ...;  
}
```

```
int *p;  
int array[10];  
  
for (p = array; p < &array[10]; p++)  
{  
    *p = ...;  
}
```

These code sequences have the same effect!

Translate Assembly

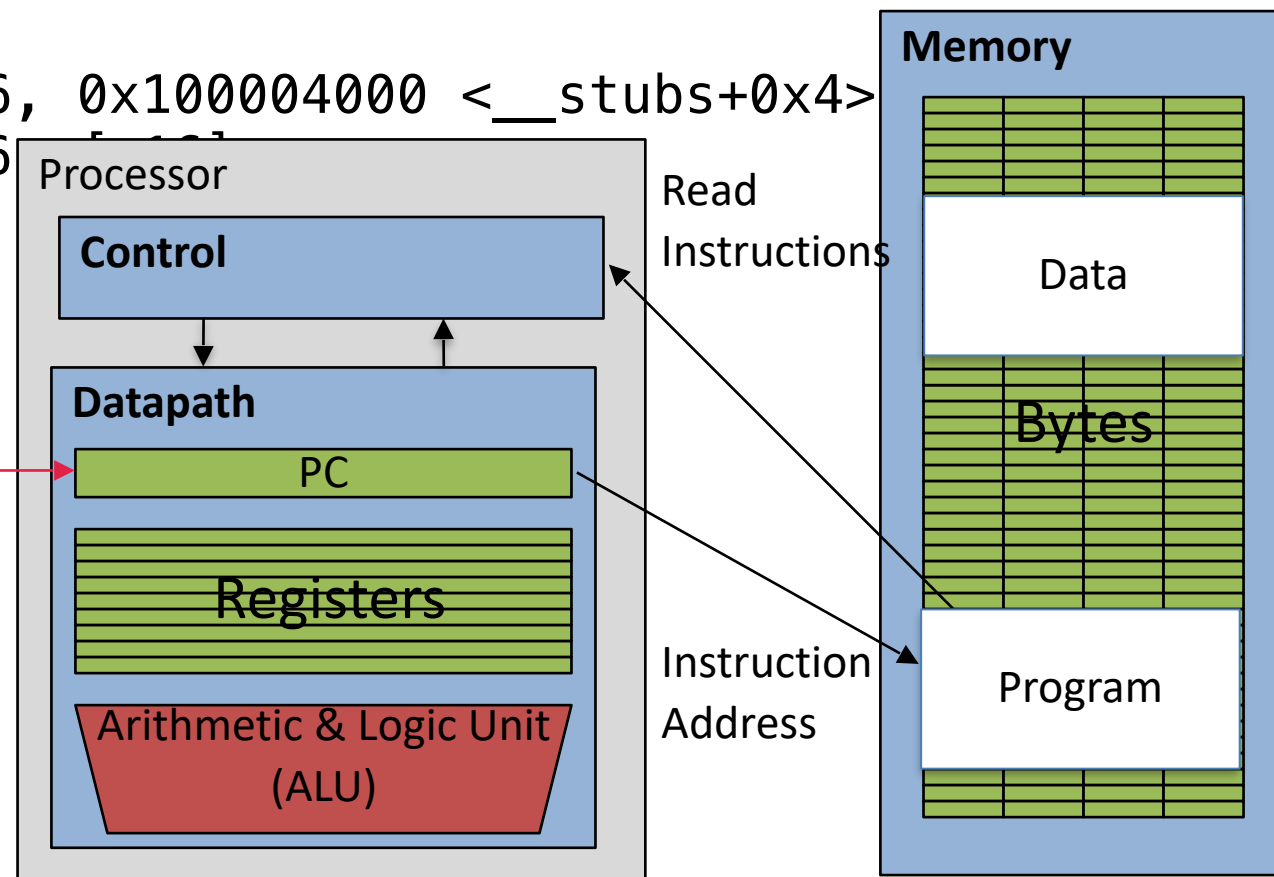
<code>addi x10, x0, 0x7</code>	<code>x10 = 7</code>
<code>add x12, x0, x0</code>	<code>x12 = 0</code>
<code>label_a:</code>	<code>label_a: x14 = x10 & 1</code>
<code>andi x14, x10, 1</code>	<code>if (x14!=0)</code>
<code>beq x14, x0, label_b</code>	<code>{x12 = x10+x12;}</code>
<code>add x12, x10, x12</code>	<code>label_b: x10 = x10-1;</code>
<code>label_b:</code>	<code>if (x10!=0)</code>
<code>addi x10, x10, -1</code>	<code>{go to label_a;}</code>
<code>bne x10, x0, label_a</code>	

Call a Function—Unconditional Jump

```
00000000100003f40 <_main>:
100003f40: ff c3 00 d1    sub sp, sp, #48
...
100003f58: 48 9a 80 52    mov w8, #1234
100003f5c: a8 83 1f b8    stur w8, [x29, #-8]
100003f60: 28 1c 82 52    mov w8, #4321
100003f64: a8 43 1f b8    stur w8, [x29, #-12]
100003f68: a8 83 5f b8    ldur w8, [x29, #-8]
100003f6c: a9 43 5f b8    ldur w9, [x29, #-12]
100003f70: 08 01 09 0b    add w8, w8, w9
...
100003f90: 05 00 00 94    bl 0x100003fa4 <_printf+0x100003fa4>
...
Disassembly of section __TEXT,__stubs:
00000000100003fa4 <__stubs>:
100003fa4: 10 00 00 b0    adrp x16, 0x100004000 <__stubs+0x4>
100003fa8: 10 02 40 f9    ldr x16, [x16, #0]
100003fac: 00 02 1f d6    br x16
```

Increase by 4
each time an
instruction
is executed

Except for
branch/jump/
function call

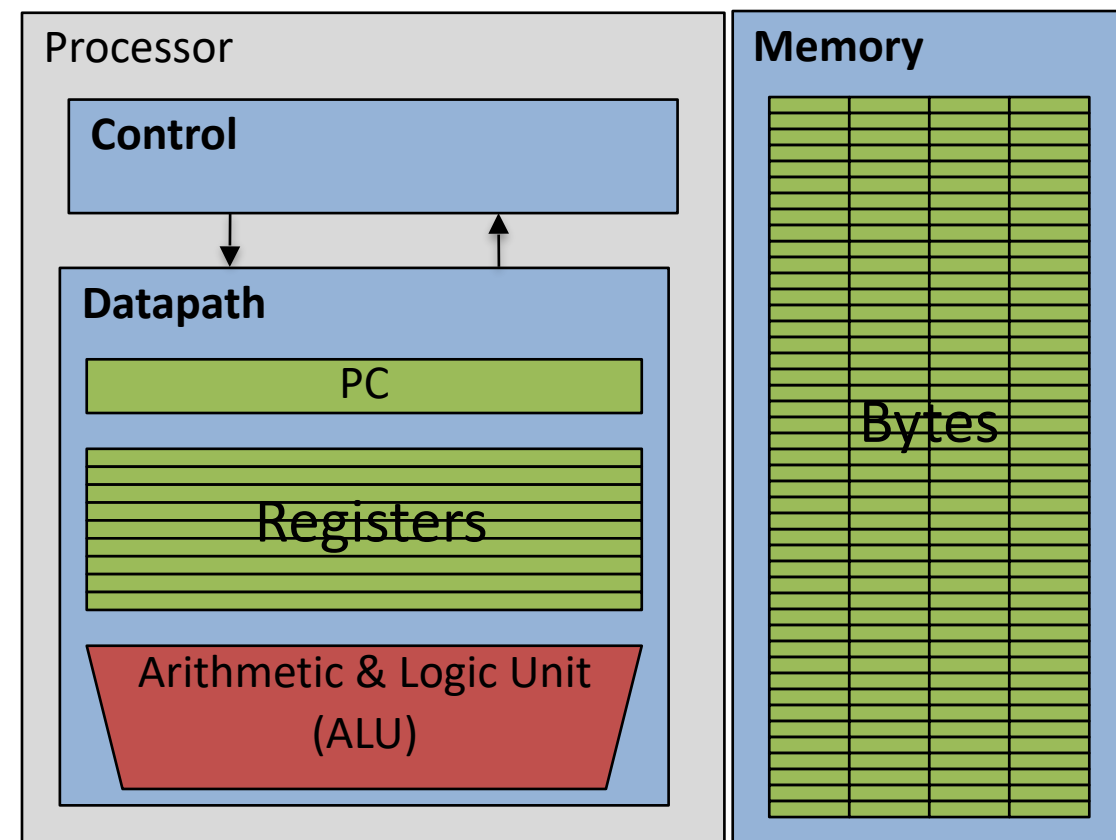


Call a Function

```
#include <stdio.h>
int sum_two_number(int a, int b)
{
    int y;
    return y=a+b;
}
int main(int argc, const char * argv[]) {
    int x=4321, y=1234;
    int a=1,b=2,c=3,d=4,e=5,f=6,g=0;
    y = sum_two_number(x,y);
    c = sum_two_number(a,b);
    f = sum_two_number(e,d);
    g = sum_two_number(c,f);
    printf("Sum is %d.\n",y);
    return 0;
}
```

1. Put parameters in a place where function can access them
2. Transfer control to function (PC jump to sum_two_number)

3. Acquire (local) storage resources needed for function
4. Perform desired task of the function

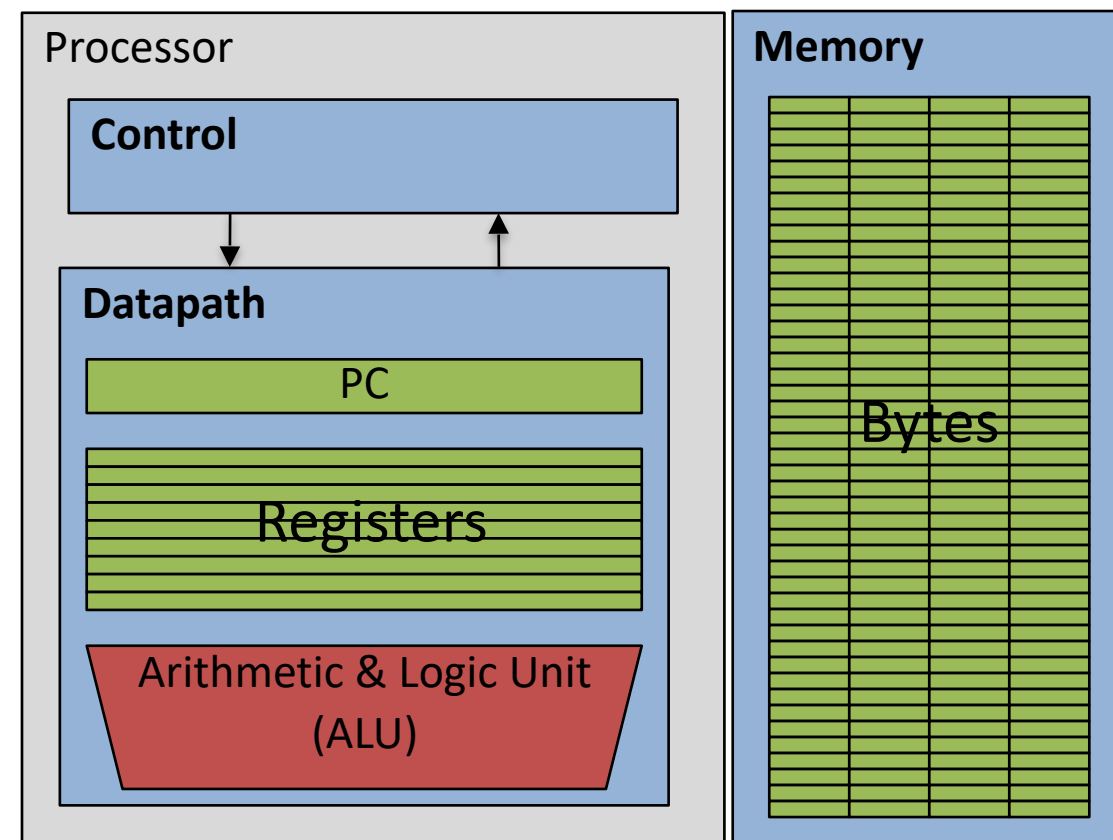


Call a Function

```
#include <stdio.h>
int sum_two_number(int a, int b)
{
    int y;
    return y=a+b;
}
int main(int argc, const char * argv[]) {
    int x=4321, y=1234;
    int a=1,b=2,c=3,d=4,e=5,f=6,g=0;
    y = sum_two_number(x,y);
    c = sum_two_number(a,b);
    f = sum_two_number(e,d);
    g = sum_two_number(c,f);
    printf("Sum is %d.\n",y);
    return 0;
}
```

5. Put result value in a place where calling code can access it and restore any registers you used

6. Return control to point of origin, since a function can be called from several points in a program



RISC-V Function Call Conventions

- Registers faster than memory, so use them as much as possible
- Give names to registers, conventions on how to use them

REGISTER NAME, USE, CALLING CONVENTION

④

REGISTER	NAME	USE	SAVER
x0	zero	The constant value 0	N.A.
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	--
x4	tp	Thread pointer	--
x5-x7	t0-t2	Temporaries	Caller
x8	s0/fp	Saved register/Frame pointer	Callee
x9	s1	Saved register	Callee
x10-x11	a0-a1	Function arguments/Return values	Caller
x12-x17	a2-a7	Function arguments	Caller
x18-x27	s2-s11	Saved registers	Callee
x28-x31	t3-t6	Temporaries	Caller

Older version: <https://riscv.org/wp-content/uploads/2015/01/riscv-calling.pdf>

Latest draft: <https://github.com/riscv-non-isa/riscv-elf-psabi-doc/releases/tag/draft-20230220-87f4a72d5aeaf048b35a230e0ba5accd1bfcf072>

RISC-V Function Call Conventions

- $a0-a7$ ($x10-x17$): eight argument registers to pass parameters and return values ($a0-a1$)
- ra : one return address register to return to the point of origin ($x1$)
- Also $s0-s1$ ($x8-x9$) and $s2-s11$ ($x18-x27$): saved registers (more about those later)

REGISTER NAME, USE, CALLING CONVENTION

④

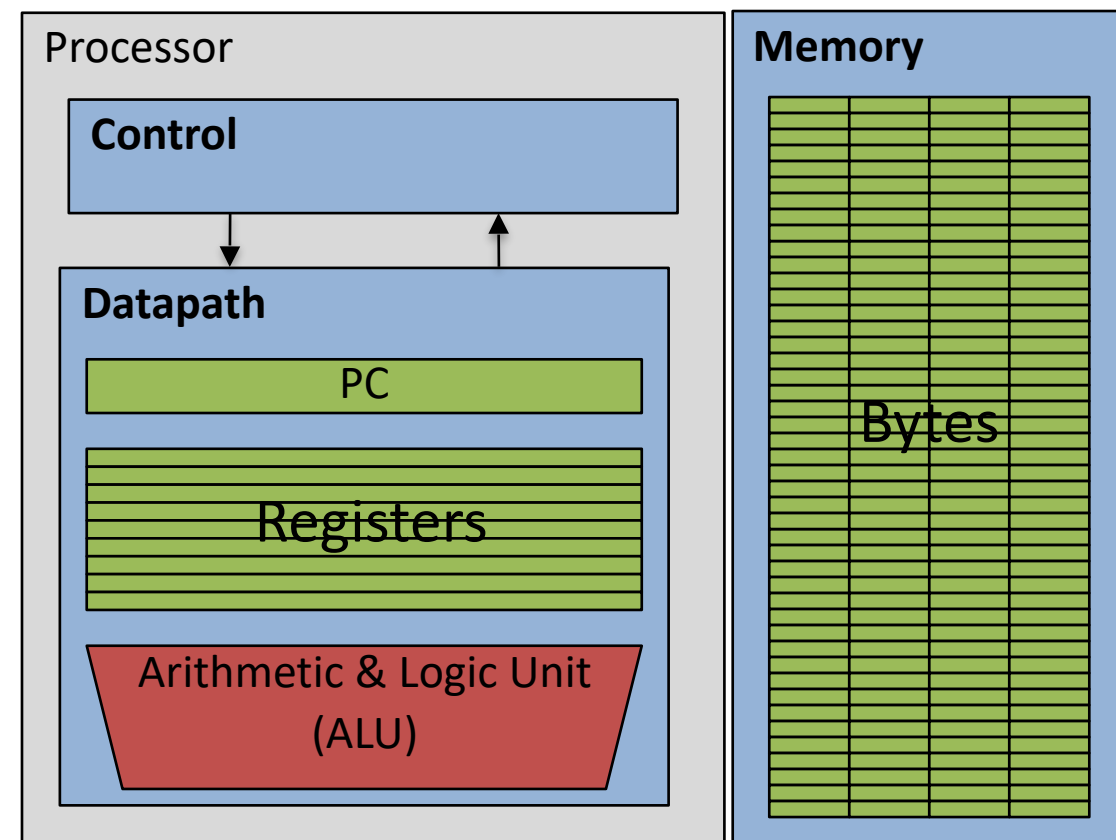
REGISTER	NAME	USE	SAVER
$x0$	zero	The constant value 0	N.A.
$x1$	ra	Return address	Caller
$x2$	sp	Stack pointer	Callee
$x3$	gp	Global pointer	--
$x4$	tp	Thread pointer	--
$x5-x7$	$t0-t2$	Temporaries	Caller
$x8$	$s0/fp$	Saved register/Frame pointer	Callee
$x9$	$s1$	Saved register	Callee
$x10-x11$	$a0-a1$	Function arguments/Return values	Caller
$x12-x17$	$a2-a7$	Function arguments	Caller
$x18-x27$	$s2-s11$	Saved registers	Callee
$x28-x31$	$t3-t6$	Temporaries	Caller

Call a Function

```
#include <stdio.h>
int sum_two_number(int a, int b)
{
    int y;
    return y=a+b;
}
int main(int argc, const char * argv[]) {
    int x=4321, y=1234;
    int a=1,b=2,c=3,d=4,e=5,f=6,g=0;
    y = sum_two_number(x,y);
    c = sum_two_number(a,b);
    f = sum_two_number(e,d);
    g = sum_two_number(c,f);
    printf("Sum is %d.\n",y);
    return 0;
}
```

y is returned function argument;
Can be put in registers a0–a1

x and y are function arguments;
Can be put in registers a0–a7



Call a Function

```
#include <stdio.h>
```

```
int sum_two_number(int a, int b)
```

```
{
```

Func_called:

0x2000 //one instruction

0x2004 //another instruction

```
}
```

... .. //**need jump back to main()**

```
int main(int argc, const char * argv[]) {
```

Start:

0x1000 //one instruction

0x1004 //another instruction

0x1008 //a third instruction

0x100c //PC jump to 0x2000 (call function
sum_two_number)

0x1010 //next instruction... ..

```
}
```

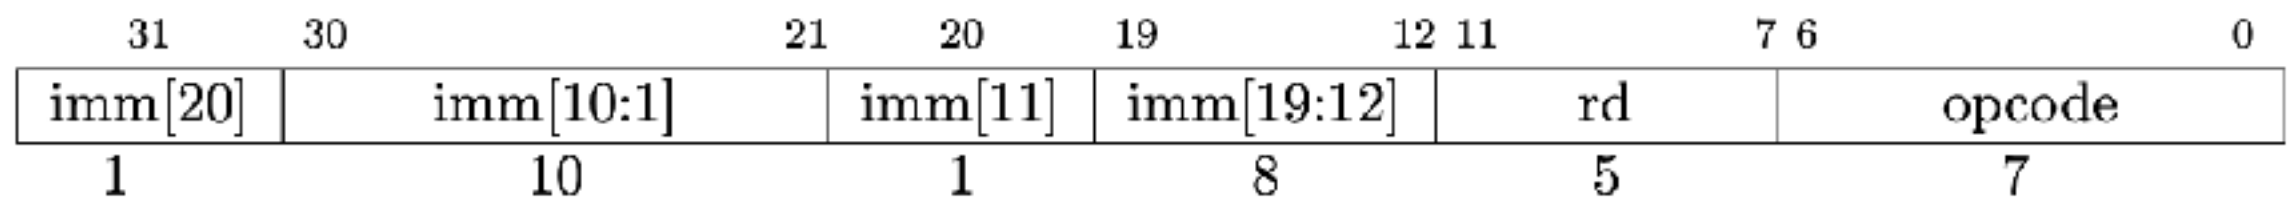
... ..

Save this value
to register ra



Call a Function—Jump

- JAL: Jump & Link, jump to function
- Unconditional jump (J-type)



`jal rd label`

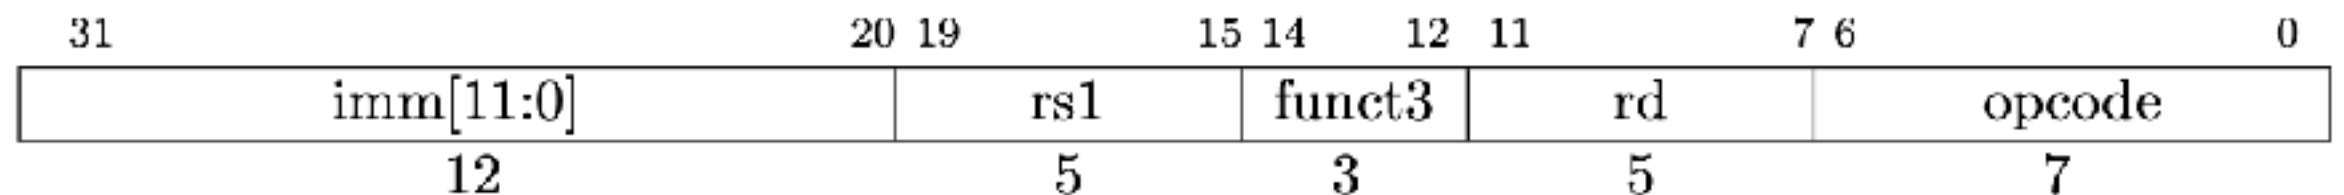
Jump to label (`imm+PC`, explain later) and save return address (`PC+4`) to `rd`;

`rd` is `x1` (`ra`) by convention; sometimes can be `x5`.

When `rd` is `x0`, it is simply unconditional jump (`j`) without recording `PC+4`.

Return—Jump

- JALR: Jump & Link Register
- Unconditional jump (**I-type**)



`jalr rd label`

Jump to label $(\text{imm} + \text{rs1}) \& \sim 1$ and save return address $(\text{PC} + 4)$ to `rd`

`rs1` can be the return address we just saved to `ra`

When `rd` is `x0`, it is simply unconditional jump (j) without recording $\text{PC} + 4$.

Jump

`-jal rd offset -jalr rd rs offset`

- Jump and Link
 - Add the immediate value to the current address in the program (the “Program Counter”), go to that location
 - The offset is 20 bits, sign extended and left-shifted one (not two)
 - At the same time, store into rd the value of PC+4
 - So we know where it came from (need to return to)
 - `jal offset == jal x1 offset` (pseudo-instruction; x1 = ra = return address)
 - `j offset == jal x0 offset` (jump is a pseudo-instruction in RISC-V)
- Two uses:
 - Unconditional jumps in loops and the like
 - Calling other functions

Jump and Link Register

- The same except the destination
 - Instead of $PC + \text{immediate}$ it is $rs + \text{immediate}$
 - Same immediate format as I-type: 12 bits, sign extended
- Again, if you don't want to record where you jump to...
 - `jr rs == jalr x0 rs`
- Two main uses
 - Returning from functions (which were called using Jump and Link)
 - Calling pointers to function

Notes on Functions

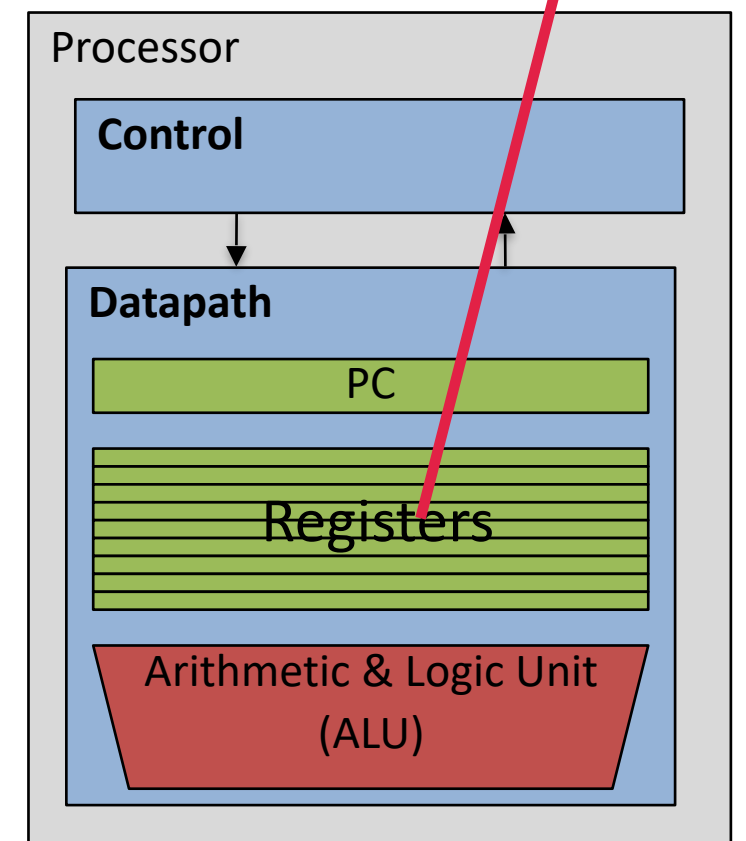
- Calling program (**caller**) puts parameters into registers **a0–a7** and uses **jal X** to invoke (**callee**) at address labeled **X**
- Must have register in computer with address of currently executing instruction
 - Instead of Instruction Address Register (better name), historically called Program Counter (PC)
 - It's a program's counter; it doesn't count programs!
- What value does **jal X** place into **ra**?
- **jr ra** puts address inside **ra** back into PC

Call a Function

1. Put parameters in a place where function can access them
2. Transfer control to function (PC jump to function)
3. Acquire (local) storage resources needed for function
4. Perform desired task of the function
5. Put result value in a place where calling code can access it and restore any registers you used
6. Return control to point of origin, since a function can be called from several points in a program

Where Are Old Register Values Saved to Restore Them After Function Call?

- Need a place to save old values before call function, restore them when return, and delete
- Ideal is stack: last-in-first-out queue (e.g., stack of plates)
 - Push: placing data onto stack
 - Pop: removing data from stack
- Stack in memory, so need register to point to it
- `sp` is the stack pointer in RISC-V (`x2`)
- Convention is grow from high to low addresses
 - Push decrements `sp`, Pop increments `sp`



Stack

- Stack frame may include:
 - Return “instruction” address
 - Parameters (spill)
 - Space for other local variables
- Stack frames contiguous; stack pointer (`sp/x2`) tells where bottom of stack frame is
- When procedure ends, stack frame is tossed off the stack; frees memory for future stack frames; `sp` restores

Example

- Leaf function: a function that calls no function

```
int Leaf (int g, int h, int i, int j)
{
    int f; f = (g + h) - (i + j);
    return f;
}
int main (void){
    int a=1, b=2, c=3, d=4, e;
    e = Leaf(a,b,d,c);
    return e;
} /*a function called by OS*/
```

0	x0/zero
ra	x1
sp	x2
... ..	
s1	x9
a0	x10
a1	x11
a2	x12
a3	x13
a4	x14
... ..	

- Parameter variables g, h, i, and j in argument registers **a1**, **a2**, **a3**, and **a4**, and f in **a0** when returned, and assume e in **s1**, later should be copied to **a0** when returned
- Register **ra** consideration

Stack Before, During, After Function

- Caller needs to save old values of a1, a2, a3 and a4
- Callee needs to save old value of s1 (and any other callee saved registers), and makes sure they are not changed after return

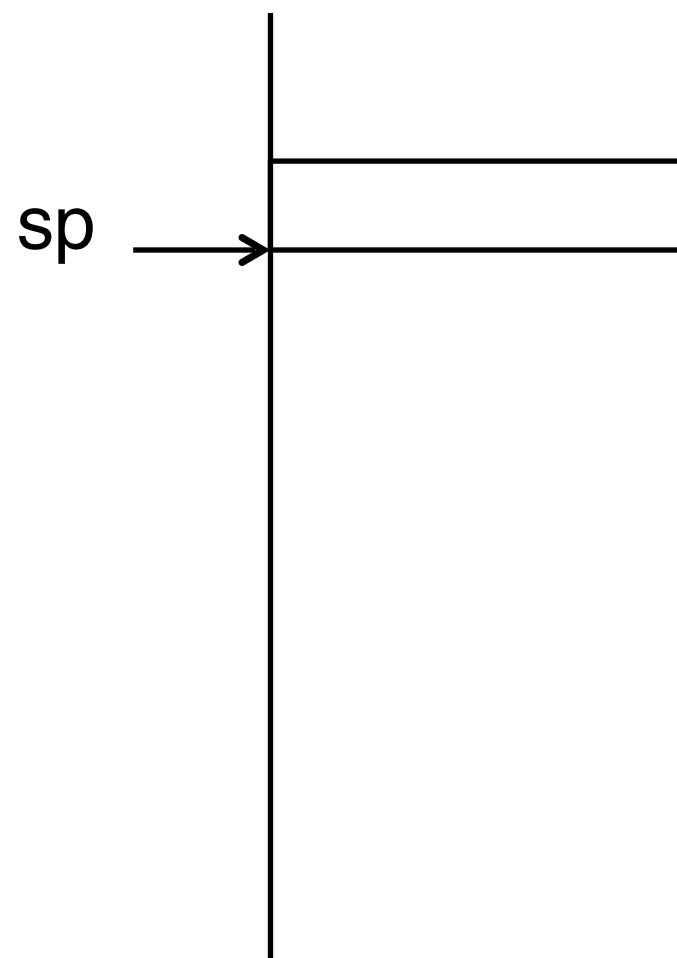
REGISTER NAME, USE, CALLING CONVENTION

④

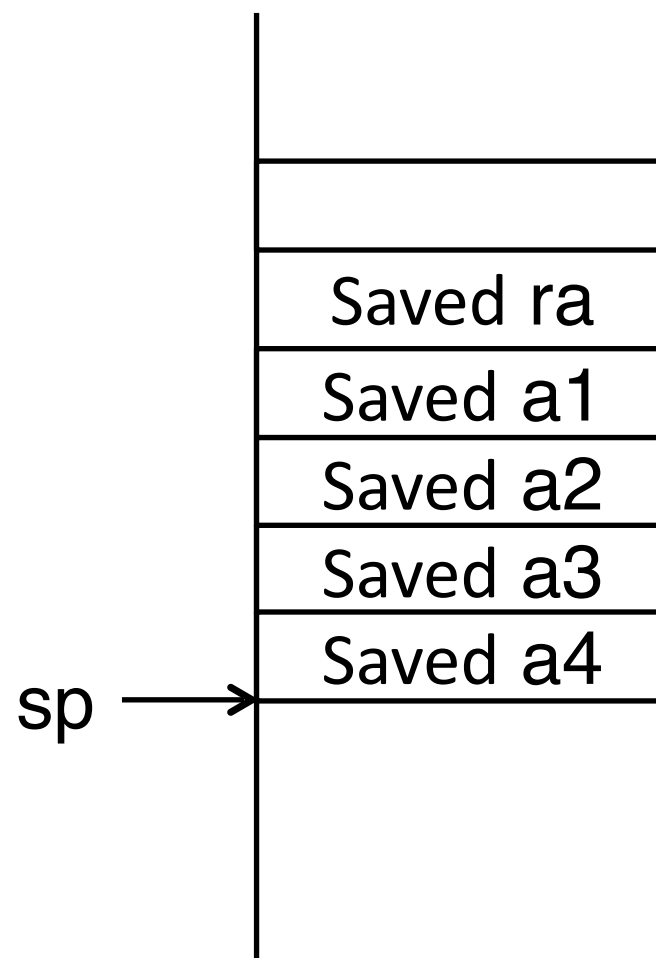
REGISTER	NAME	USE	SAVER
x0	zero	The constant value 0	N.A.
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	--
x4	tp	Thread pointer	--
x5-x7	t0-t2	Temporaries	Caller
x8	s0/fp	Saved register/Frame pointer	Callee
x9	s1	Saved register	Callee
x10-x11	a0-a1	Function arguments/Return values	Caller
x12-x17	a2-a7	Function arguments	Caller
x18-x27	s2-s11	Saved registers	Callee
x28-x31	t3-t6	Temporaries	Caller

Stack Before, During, After Function

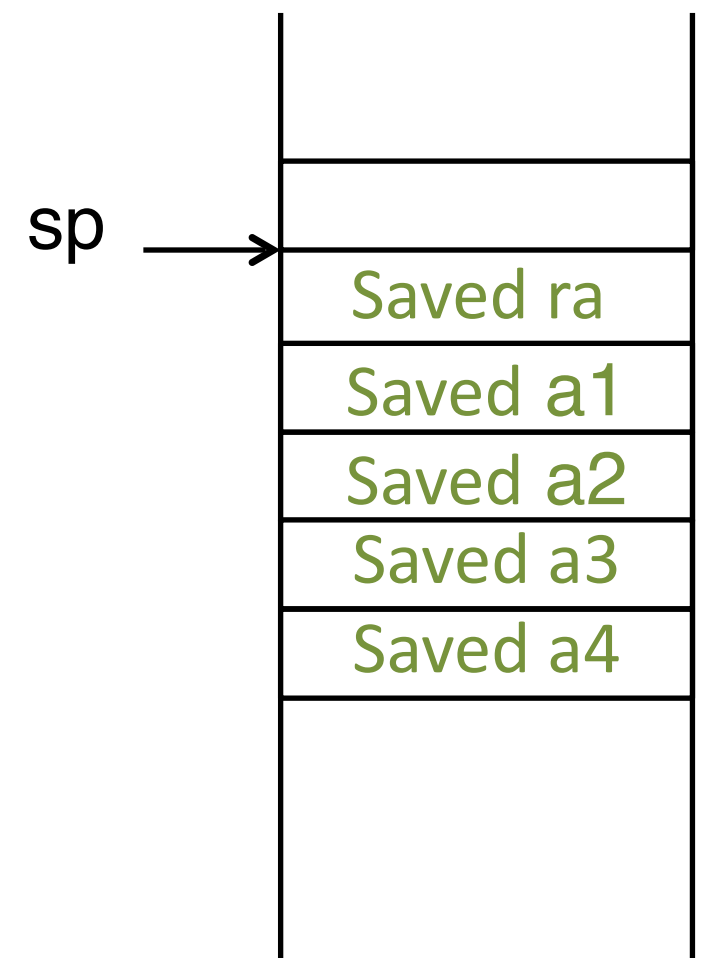
- Need to save old values of ra, a1, a2, a3 and a4 (caller-saved)
- W.r.t. main()



Before call



During call



After call

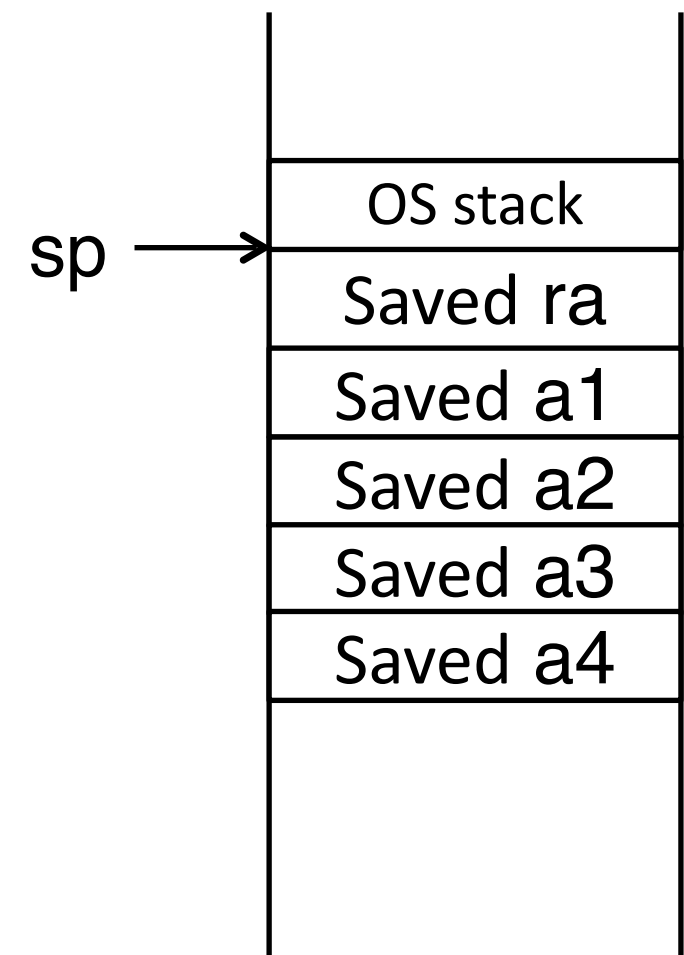
RISC-V Code for Main()/Leaf()

Main:

```
addi    sp, sp, -20  # adjust stack for 5 items, 4 int & 1 ra pointer/address
sw      ra, 16(sp)   # save ra for use afterwards (return to OS)
sw      a1, 12(sp)   # save a1 for use afterwards, these are all caller-saved
sw      a2, 8(sp)    # save a2 for use afterwards
sw      a3, 4(sp)    # save a3 for use afterwards
sw      a4, 0(sp)    # save a4 for use afterwards

jal     ra, Leaf     # save a1 for use afterwards

lw      a1, 12(sp)   # restore register a1
lw      a2, 8(sp)    # restore register a2
lw      a3, 8(sp)    # restore register a2
lw      a4, 8(sp)    # restore register a2
lw      ra, 8(sp)    # restore register ra
addi    sp, sp, 20   # adjust stack to delete 5 items
mv      a0, a0       # move result to return register
jr      ra           # return
```



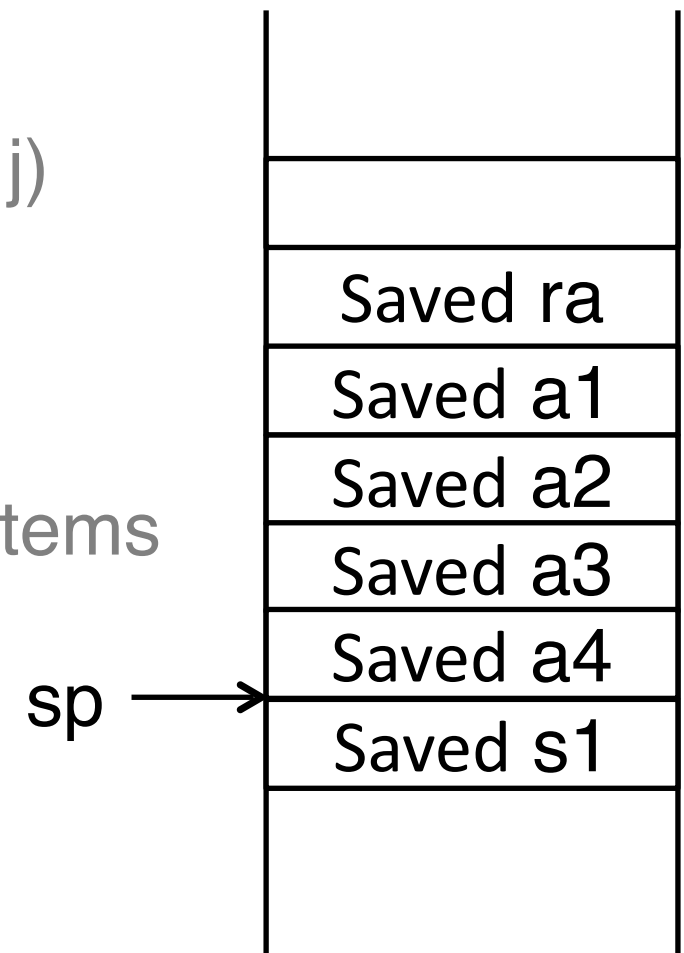
Stack
During call

RISC-V Code for Main()/Leaf()

Leaf:

```
addi sp, sp, -4    # adjust stack for 1 items, callee saved s1
sw   s1, 0(sp)     # save callee saved s1 to stack
add  a1, a0, a1     # g = g + h
add  a2, a2, a3     # j = i + j
sub  s1, a0, a2     # calculate result (g + h) - (i + j)
mv   a0, s1        # return value (g + h) - (i + j)

lw   s1, 0(sp)     # restore register s1 for caller
addi sp, sp, 4     # adjust stack to delete 1 items
jr   ra           # jump back to caller
                        (pseudo-assembly: ret)
```



Stack
During call

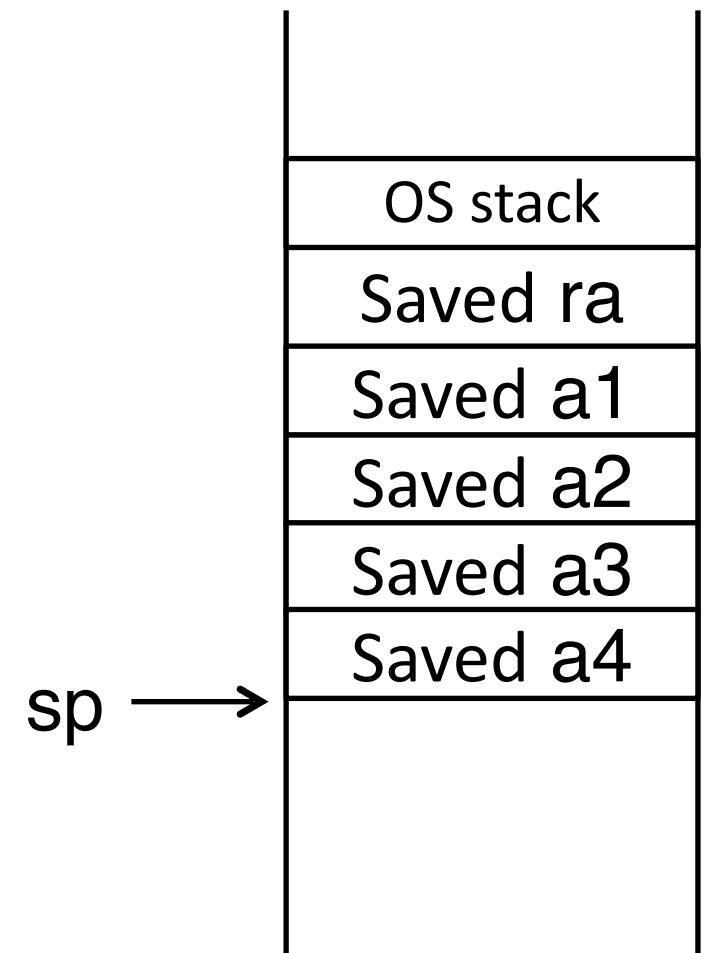
RISC-V Code for Main()/Leaf()

Main:

```
addi    sp, sp, -20  # adjust stack for 5 items, 4 int & 1 ra pointer/address
sw      ra, 16(sp)   # save ra for use afterwards (return to OS)
sw      a1, 12(sp)   # save a1 for use afterwards, these are all caller-saved
sw      a2, 8(sp)    # save a2 for use afterwards
sw      a3, 4(sp)    # save a3 for use afterwards
sw      a4, 0(sp)    # save a4 for use afterwards

jal     ra, Leaf     # save a1 for use afterwards

lw      a1, 12(sp)   # restore register a1
lw      a2, 8(sp)    # restore register a2
lw      a3, 8(sp)    # restore register a2
lw      a4, 8(sp)    # restore register a2
lw      ra, 8(sp)    # restore register ra
addi    sp, sp, 20   # adjust stack to delete 5 items
mv      a0, a0       # move result to return register
jr      ra           # return
```



Stack
During call

Call a Function

1. Caller put parameters in a place where function can access them (a0-a7, or stack when registers not avail.), and then save caller-saved registers to stack
2. Transfer control to function (PC jump to function): JAL, ra is changed to where caller left
3. Acquire (local) storage resources needed for function: change sp (size decided when compiling);
Push callee-saved registers to stack (e.g., s0-s11)
4. Perform desired task of the function
5. Put result value in a place where calling code can access it (a0, a1), and restore callee-saved registers (s0-s11, sp)
6. Return control to point of origin, since a function can be called from several points in a program (jr); caller restores caller-saved registers