

# CS 110

# Computer Architecture

## *Dependability and RAID*

Instructors:

**Chundong Wang & Siting Liu**

<https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html>

**School of Information Science and Technology SIST**

**ShanghaiTech University**

**Slides based on UC Berkley's CS61C**

# Review

- DMA
- Disk I/Os
- Flash memory
- Network
  - Abstraction
  - Protocols


# A Story of Dependability



Key locked in flat



Spare key in office



Redundancy

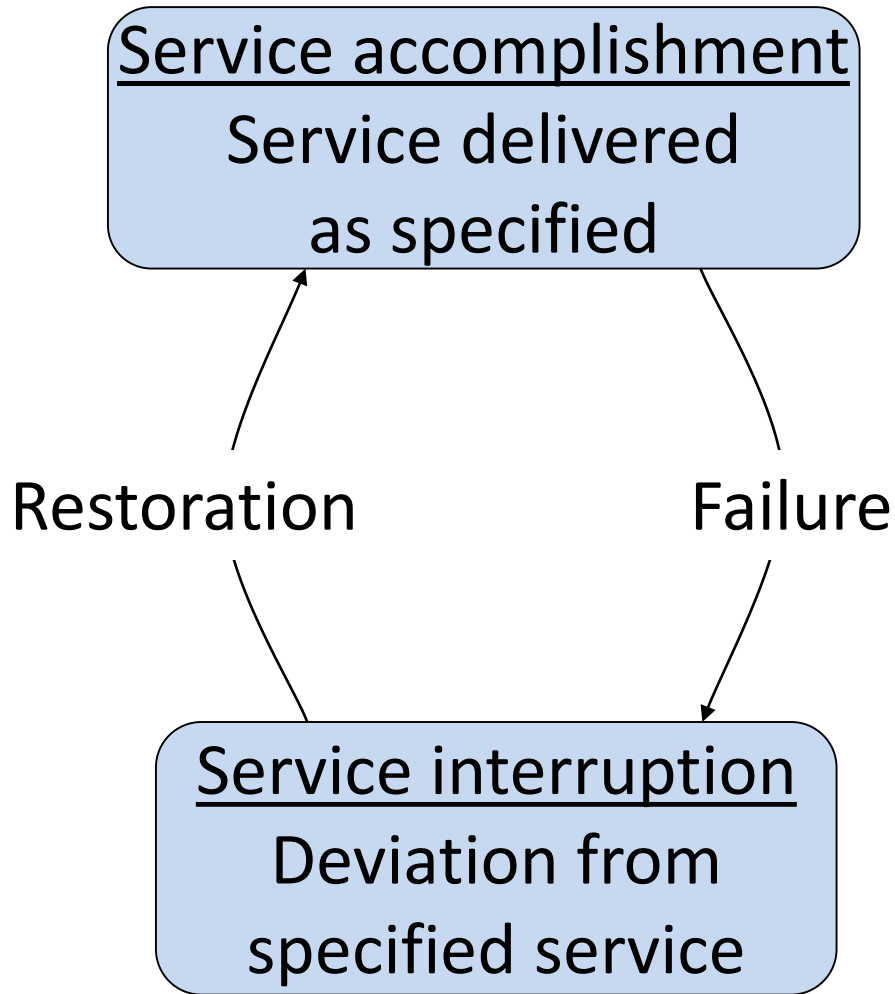
# Great Idea #6:

## Dependability via Redundancy

- Applies to everything from datacenters to memory
  - Redundant datacenters so that can lose 1 datacenter but Internet service stays online
  - Redundant routes so can lose nodes but Internet doesn't fail
  - Redundant memory bits of so that can lose 1 bit but no data (Error Correcting Code/ECC Memory)
  - Redundant disks so that can lose 1 disk but not lose data (Redundant Arrays of Independent Disks/RAID)



# Dependability



- Fault: failure of a component
  - May or may not lead to system failure

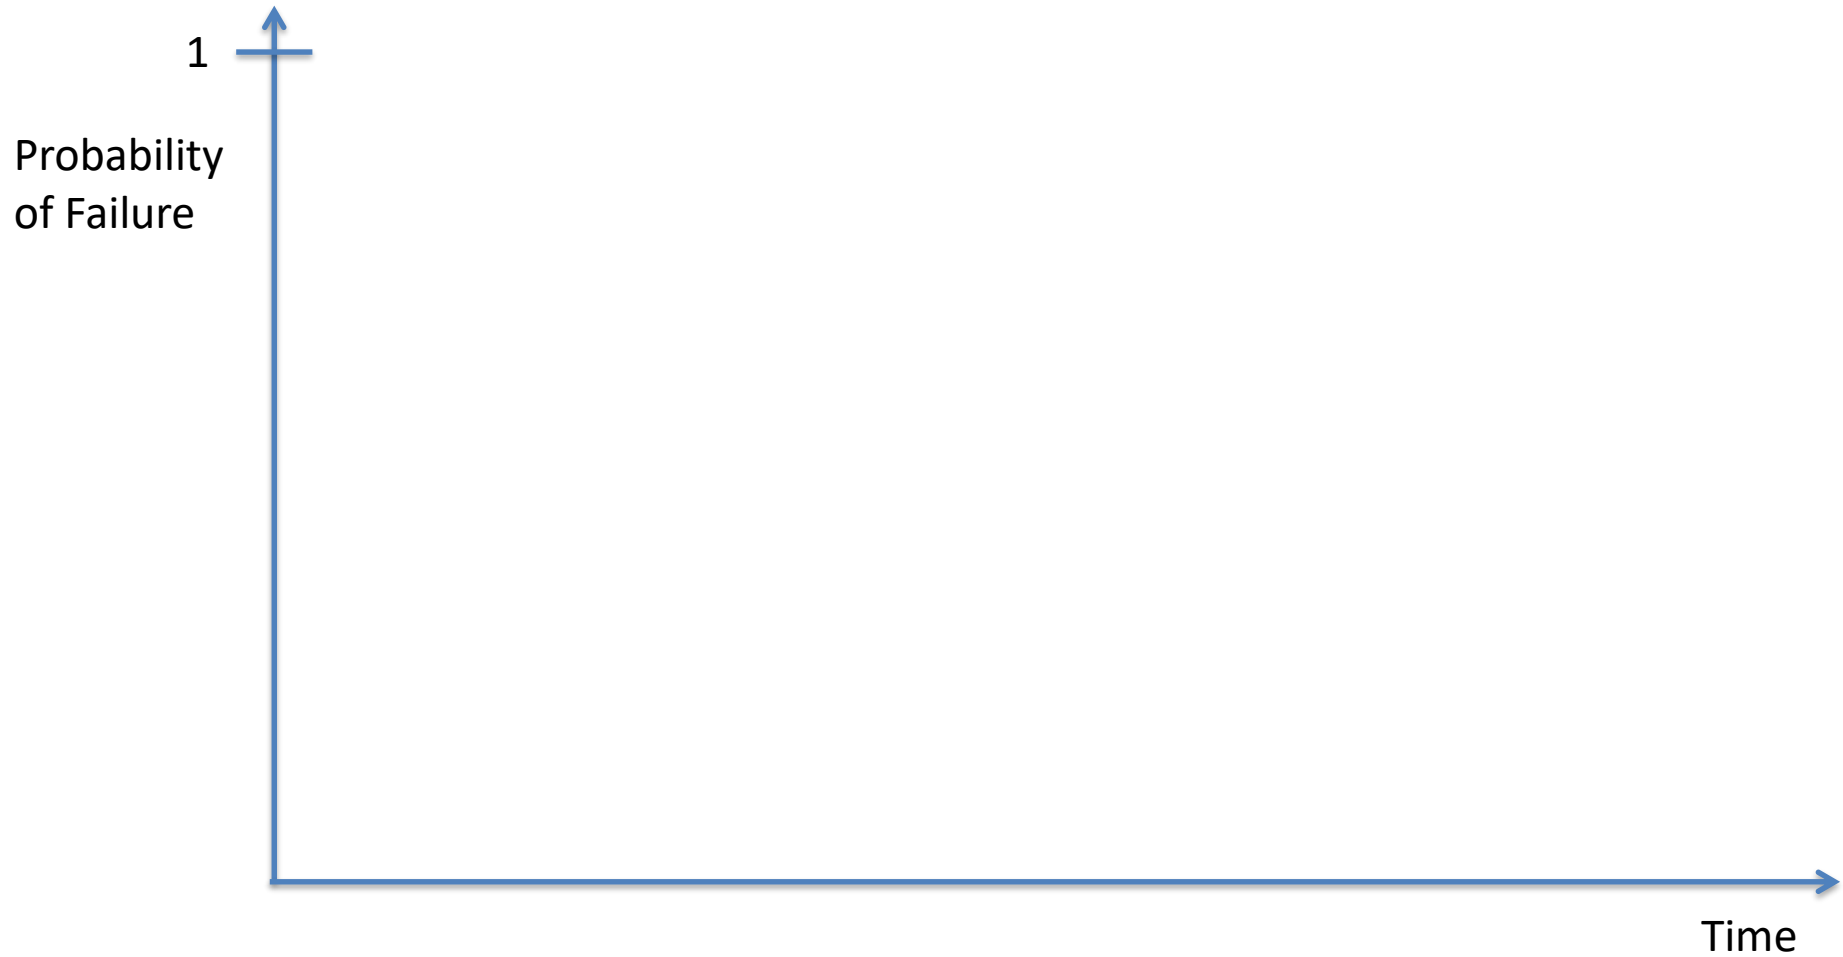
# Dependability via Redundancy: Time vs. Space

- *Spatial Redundancy* – replicated data or check information or hardware to handle **hard** and soft (transient) failures
- *Temporal Redundancy* – redundancy in time (retry) to handle soft (transient) failures

# Dependability Measures


- Reliability: Mean Time To Failure (**MTTF**)
- Service interruption: Mean Time To Repair (**MTTR**)
- Mean time between failures (**MTBF**)
  - $MTBF = MTTF + MTTR$
- Availability =  $\frac{MTTF}{MTTF + MTTR}$
- Improving Availability
  - Increase MTTF: More reliable hardware/software + Fault Tolerance
  - Reduce MTTR: improved tools and processes for diagnosis and repair

# Understanding MTTF





# Understanding MTTF



# Availability Measures

- Availability =  $\frac{MTTF}{MTTF+MTTR}$  as %
  - MTTF, MTBF usually measured in hours
- Since hope rarely down, shorthand is “number of 9s of availability per year”
- 1 nine: 90% => 36 days of repair/year
- 2 nines: 99% => 3.6 days of repair/year
- 3 nines: 99.9% => 526 minutes of repair/year
- 4 nines: 99.99% => 53 minutes of repair/year
- 5 nines: 99.999% => 5 minutes of repair/year

# Reliability Measures

- Another is average number of failures per year:  
**Annualized Failure Rate (AFR)**
  - E.g., 1000 disks with 100,000 hours MTTF
  - 365 days \* 24 hours = 8760 hours
  - $(1000 \text{ disks} * 8760 \text{ hrs/year}) / 100,000 = 87.6$  failed disks per year on average
  - $87.6 / 1000 = 8.76\%$  annual failure rate
- Google's 2007 study\* found that actual AFRs for individual drives ranged from 1.7% for first year drives to over 8.6% for three-year old drives

*\*research.google.com/archive/disk\_failures.pdf*

# Dependability Design Principle

- Design Principle: No single points of failure
  - “Chain is only as strong as its weakest link”
  - *Achilles' Heel*
- Dependability Corollary of Amdahl's Law
  - Doesn't matter how dependable you make one portion of system
  - Dependability limited by part you do not improve

# Error Detection/Correction Codes

- Memory systems generate errors (accidentally flipped-bits)
  - DRAMs store very little charge per bit
  - “Soft” errors occur occasionally when cells are struck by alpha particles or other environmental upsets
  - “Hard” errors can occur when chips permanently fail
  - Problem gets worse as memories get denser and larger
- Memories protected against failures with EDC/ECC
- Extra bits are added to each data-word
  - Used to detect and/or correct faults in the memory system
  - Each data word value mapped to unique *code word*
  - A fault changes valid code word to invalid one, which can be detected

# Block Code Principles

- Hamming distance = difference in # of bits
- $p = 0\underline{1}1\underline{0}11$ ,  $q = 0\underline{0}1\underline{1}11$ , Ham. distance  $(p,q) = 2$
- $p = 011011$ ,  
 $q = 110001$ ,  
distance  $(p,q) = ?$
- Can think of extra bits as creating a code with the data
  - There is Ham. distance between codes



Richard Hamming, 1915-98  
Turing Award Winner

# Parity



- Parity bits are added to a word to make it
  - either odd: odd numbers of '1'
  - or even: even number of '1'
- Let us add one parity bit to three-bit word

Odd Parity		Even Parity	
000	0001	000	0000
100	1000	100	1001
101	1011	101	1010
111	1110	111	1111

# Parity: Simple Error-Detection Coding

- Each data value, before it is written to memory is “tagged” with an extra bit to force the stored word to have *even parity*:
- Each word, as it is read from memory is “checked” by finding its parity (including the parity bit).

*parity*:



- A non-zero parity check indicates an error occurred:
  - 2 errors (on different bits) are not detected
  - nor any even number of errors, just odd numbers of errors are detected
- Minimum Hamming distance of valid parity codes is 2



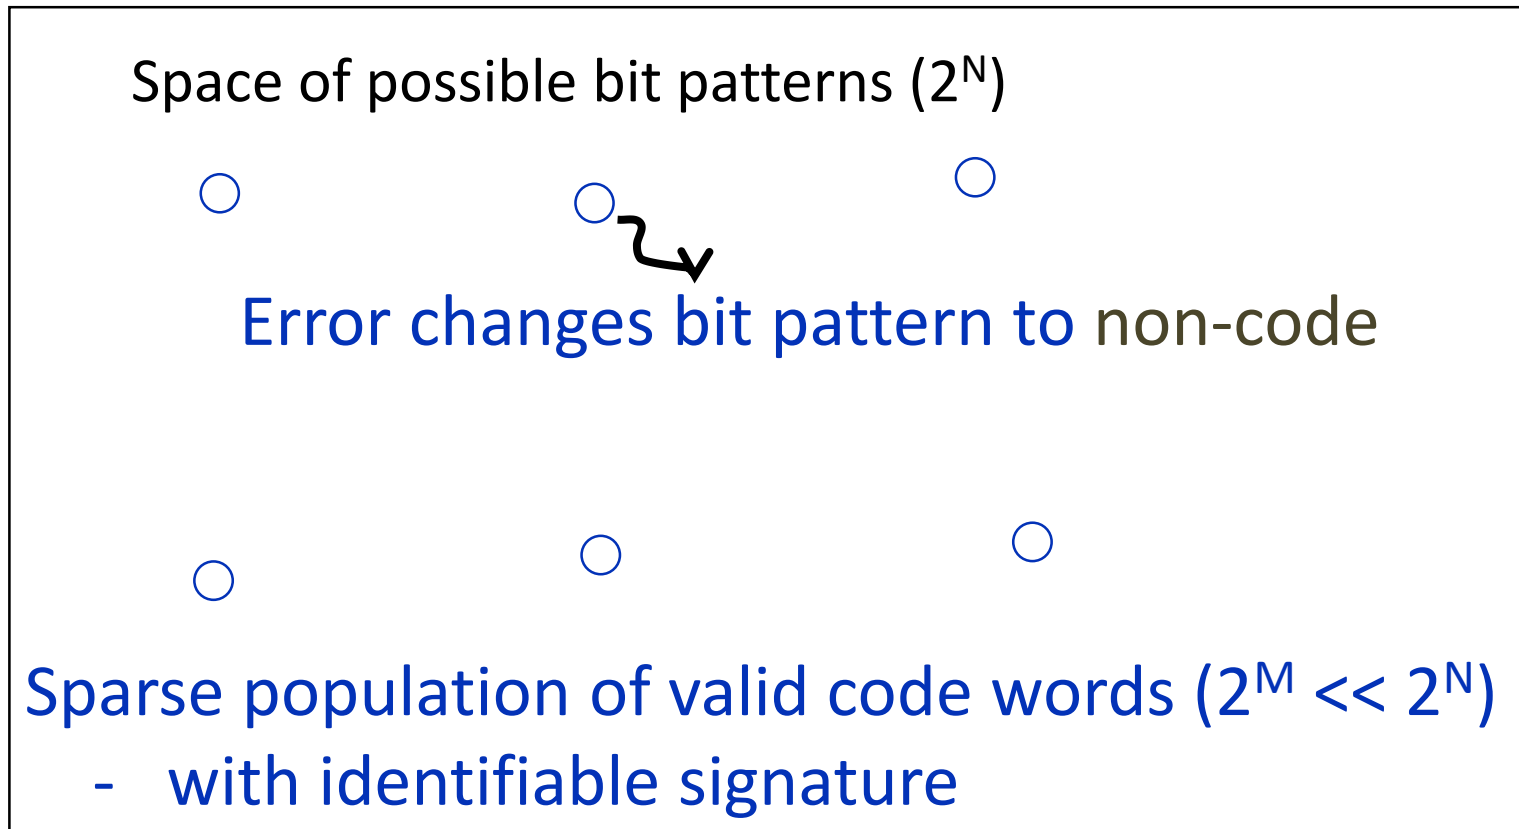
# Parity Example

- Data 0101 0101
- 4 ones, even parity now
- Write to memory:  
0101 0101 0  
to keep parity even
- Data 0101 0111
- 5 ones, odd parity now
- Write to memory:  
0101 0111 1  
to make parity even
- Read from memory  
0101 0101 0
- 4 ones => even parity,  
so no error
- Read from memory  
1101 0101 0
- 5 ones => odd parity,  
so error
- What if error in parity  
bit?

# Suppose Want to Correct 1 Error?


- Richard Hamming came up with simple to understand mapping to allow Error Correction at minimum distance of 3
  - Single error correction, double error detection
- Called “Hamming ECC”
  - Worked weekends on relay computer with unreliable card reader, frustrated with manual restarting
  - Got interested in error correction; published 1950
  - R. W. Hamming, “Error Detecting and Correcting Codes,” *The Bell System Technical Journal*, Vol. XXVI, No 2 (April 1950) pp 147-160.

# Detecting/Correcting Code Concept




- **Detection**: bit pattern fails codeword check
- **Correction**: map to nearest valid code word

# Hamming Distance: 8 code words



# Hamming Distance 2: Detection


## *Detect Single Bit Errors*



- No 1 bit error goes to another valid codeword
- $\frac{1}{2}$  codewords are valid

# Hamming Distance 3: Correction

Correct Single Bit Errors, Detect Double Bit Errors



- No 2 bit error goes to another valid codeword; 1 bit error near
- 1/4 codewords are valid

# Hamming Error Correction Code

- Use of **extra parity bits** to allow the position identification of a single error
  1. Mark all bit positions that are **powers of 2** as parity bits (positions 1, 2, 4, 8, 16, ...)
    - Start numbering bits at 1 at left (not at 0 on right)
  2. All **other bit positions** are data bits (positions 3, 5, 6, 7, 9, 10, 11, 12, 13, 14, 15, ...)
  3. Each data bit is covered by 2 or more parity bits

# Hamming ECC

4. The **position of parity bit** determines sequence of data bits that it checks
- **Bit 1 ( $0001_2$ )**: checks bits (1,3,5,7,9,11,...)
    - Bits with least significant bit of address = 1
  - **Bit 2 ( $0010_2$ )**: checks bits (2,3,6,7,10,11,14,15,...)
    - Bits with 2<sup>nd</sup> least significant bit of address = 1
  - **Bit 4 ( $0100_2$ )**: checks bits (4-7, 12-15, 20-23, ...)
    - Bits with 3<sup>rd</sup> least significant bit of address = 1
  - **Bit 8 ( $1000_2$ )**: checks bits (8-15, 24-31, 40-47 ,...)
    - Bits with 4<sup>th</sup> least significant bit of address = 1



# Graphic of Hamming Code

Bit position		1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11
Parity bit coverage	p1	X		X		X		X		X		X		X		X
	p2		X	X			X	X			X	X			X	X
	p4				X	X	X	X					X	X	X	X
	p8								X	X	X	X	X	X	X	X

- [http://en.wikipedia.org/wiki/Hamming\\_code](http://en.wikipedia.org/wiki/Hamming_code)

# Hamming ECC

5. Set parity bits to create **even parity** for each group
- A byte of data: 10011010
  - Create the coded word, leaving spaces for the parity bits:
  - |   |   |   |   |   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|---|---|---|---|
| _ | _ | 1 | _ | 0 | 0 | 1 | _ | 1 | 0 | 1 | 0 |
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | A | B | C |
  - Calculate the parity bits

# Hamming ECC

\_ \_ 1 \_ 0 0 1 \_ 1 0 1 0

- Position 1 checks bits 1, 3, 5, 7, 9, 11:

? \_ 1 \_ 0 0 1 \_ 1 0 1 0. set position 1:

**0** \_ 1 \_ 0 0 1 \_ 1 0 1 0

- Position 2 checks bits 2, 3, 6, 7, 10, 11:

0 ? 1 \_ 0 0 1 \_ 1 0 1 0. set position 2:

0 **1** 1 \_ 0 0 1 \_ 1 0 1 0

- Position 4 checks bits 4, 5, 6, 7, 12:

0 1 1 ? 0 0 1 \_ 1 0 1 0. set position 4:

0 1 1 **1** 0 0 1 \_ 1 0 1 0

- Position 8 checks bits 8, 9, 10, 11, 12:

– 0 1 1 1 0 0 1 ? 1 0 1 0. set position 8:

– 0 1 1 1 0 0 1 **0** 1 0 1 0

# Hamming ECC

- **Final** code word: 01100101010
- Data word:                   1   001   1010

# Hamming ECC Error Check

- Suppose receive

011100101110

0 1 1 1 0 0 1 0 1 1 1 0

Bit position		1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11
Parity bit coverage	p1	X		X		X		X		X		X		X		X
	p2		X	X			X	X			X	X			X	X
	p4				X	X	X	X					X	X	X	X
	p8								X	X	X	X	X	X	X	X

# Hamming ECC Error Check

- Suppose receive

011100101110

0 1 0 1 1 1    ✓

11    01    11    X-Parity 2 in error

1001    0    ✓

01110    X-Parity 8 in error

- *Implies position 8+2=10 is in error*

011100101**1**10

# Hamming ECC Error Correct

- Flip the incorrect bit ...

011100101010

- Double check

011100101010

0 1 0 1 1 1 √

11 01 01 √

1001 0 √

01 010 √

# Hamming ECC Error Detect

- Suppose receive

010100001010  
0 0 0 0 1 1    ✓  
10    00    01    ✓  
      1000        0    X  
          01010    ✓

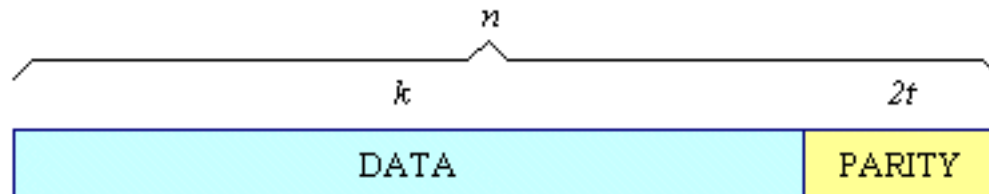
Two errors can be detected,  
but not correctable

How about  $\geq 3$  bits error?



# Cyclic Redundancy Check

- Parity is not powerful enough to detect long runs of errors (also known as *burst errors*)
- Better Alternative: *Reed-Solomon Codes*
  - Used widely in CDs, DVDs, Magnetic Disks
  - RS(255,223) with 8-bit symbols: each codeword contains 255 code word bytes (223 bytes are data and 32 bytes are parity)



- For this code:  $n = 255$ ,  $k = 223$ ,  $s = 8$ ,  $2t = 32$ ,  $t = 16$
- Decoder can correct any errors in up to 16 bytes anywhere in the codeword

# RAID: Redundancy for Disks


- Why we still worry about disks?
  - Trade-off: price, capacity, density, etc.
  - When you need storage space in petabytes (PB) or exabytes (EB)
    - 1 PB = 1024 TB
    - 1 EB = 1024 PB
  - Do not forget that flash-based SSDs also fail
    - Limited program/erase cycles ← wear levelling

# RAID: Redundant Arrays of (Inexpensive) Disks


- Files are “striped” across multiple disks
- Redundancy yields high data availability
  - Availability: service still provided to user, even if some components failed
- Disks will still fail
- Contents reconstructed from data redundantly stored in the array
  - ➔ Capacity penalty to store redundant info
  - ➔ Bandwidth penalty to update redundant info

# RAID 0: Striping

- RAID 0 provides **no** fault tolerance or redundancy
  - Striping, or disk spanning
  - High performance

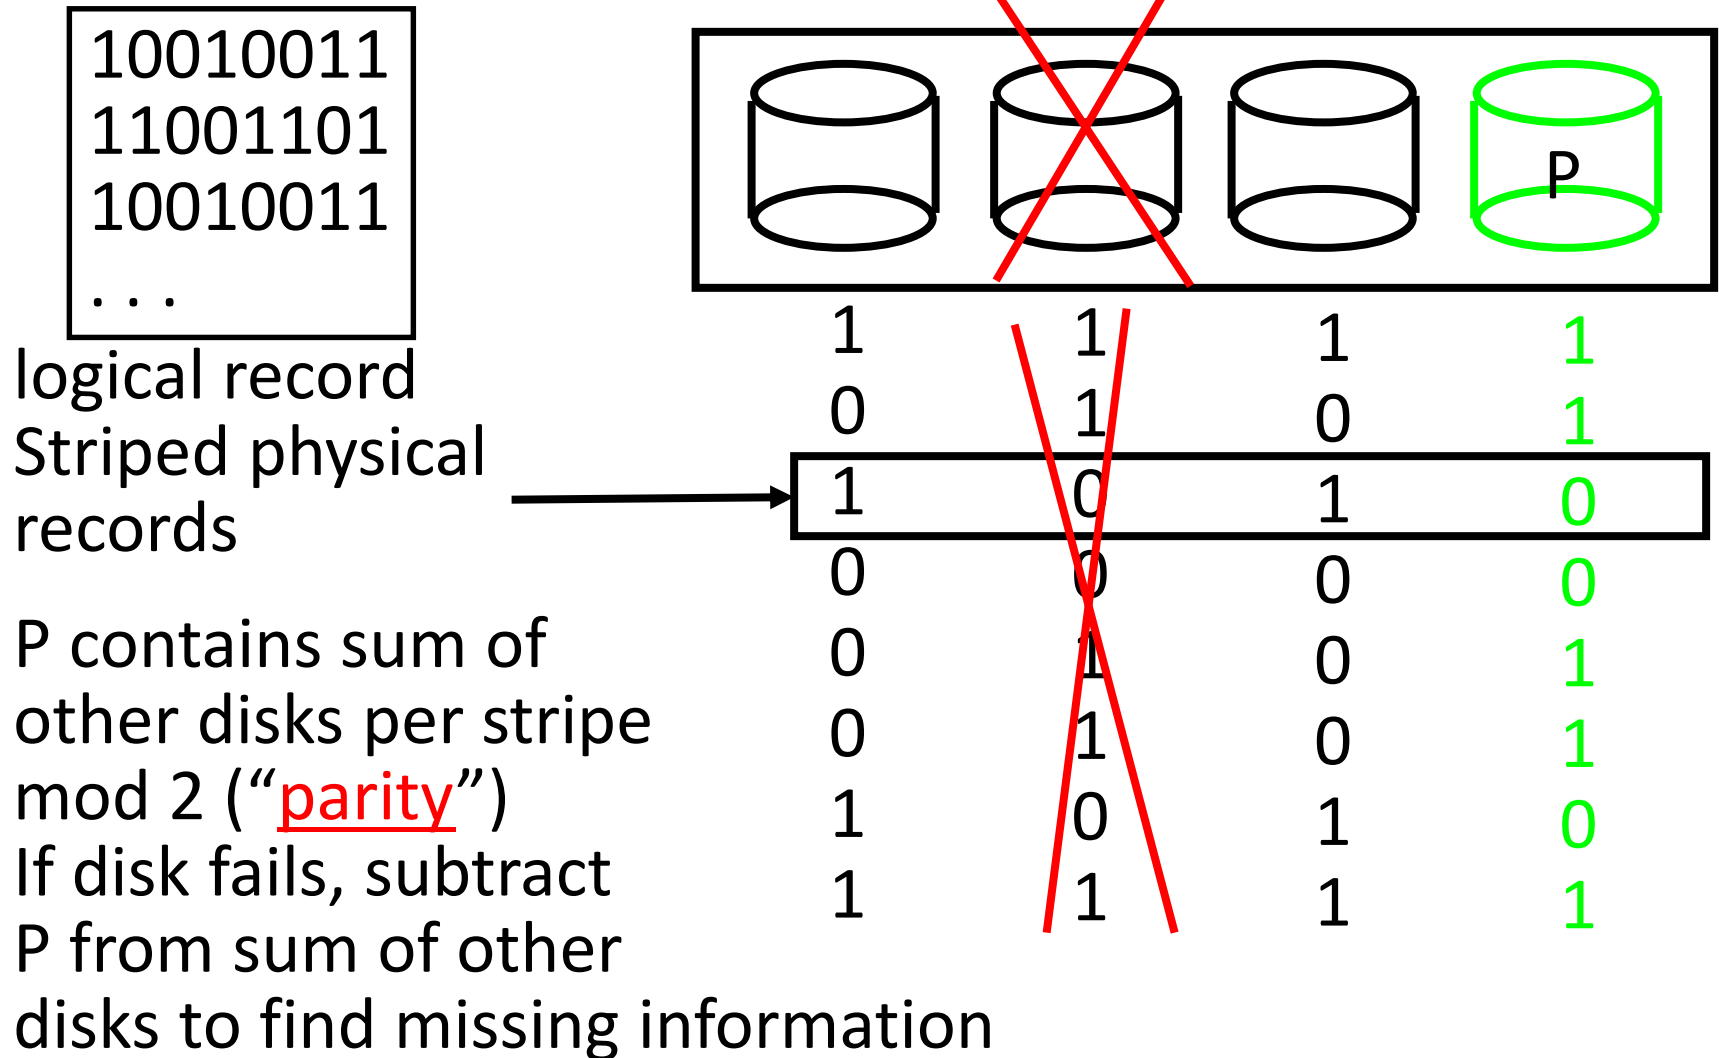


# RAID 1: Disk Mirroring/Shadowing

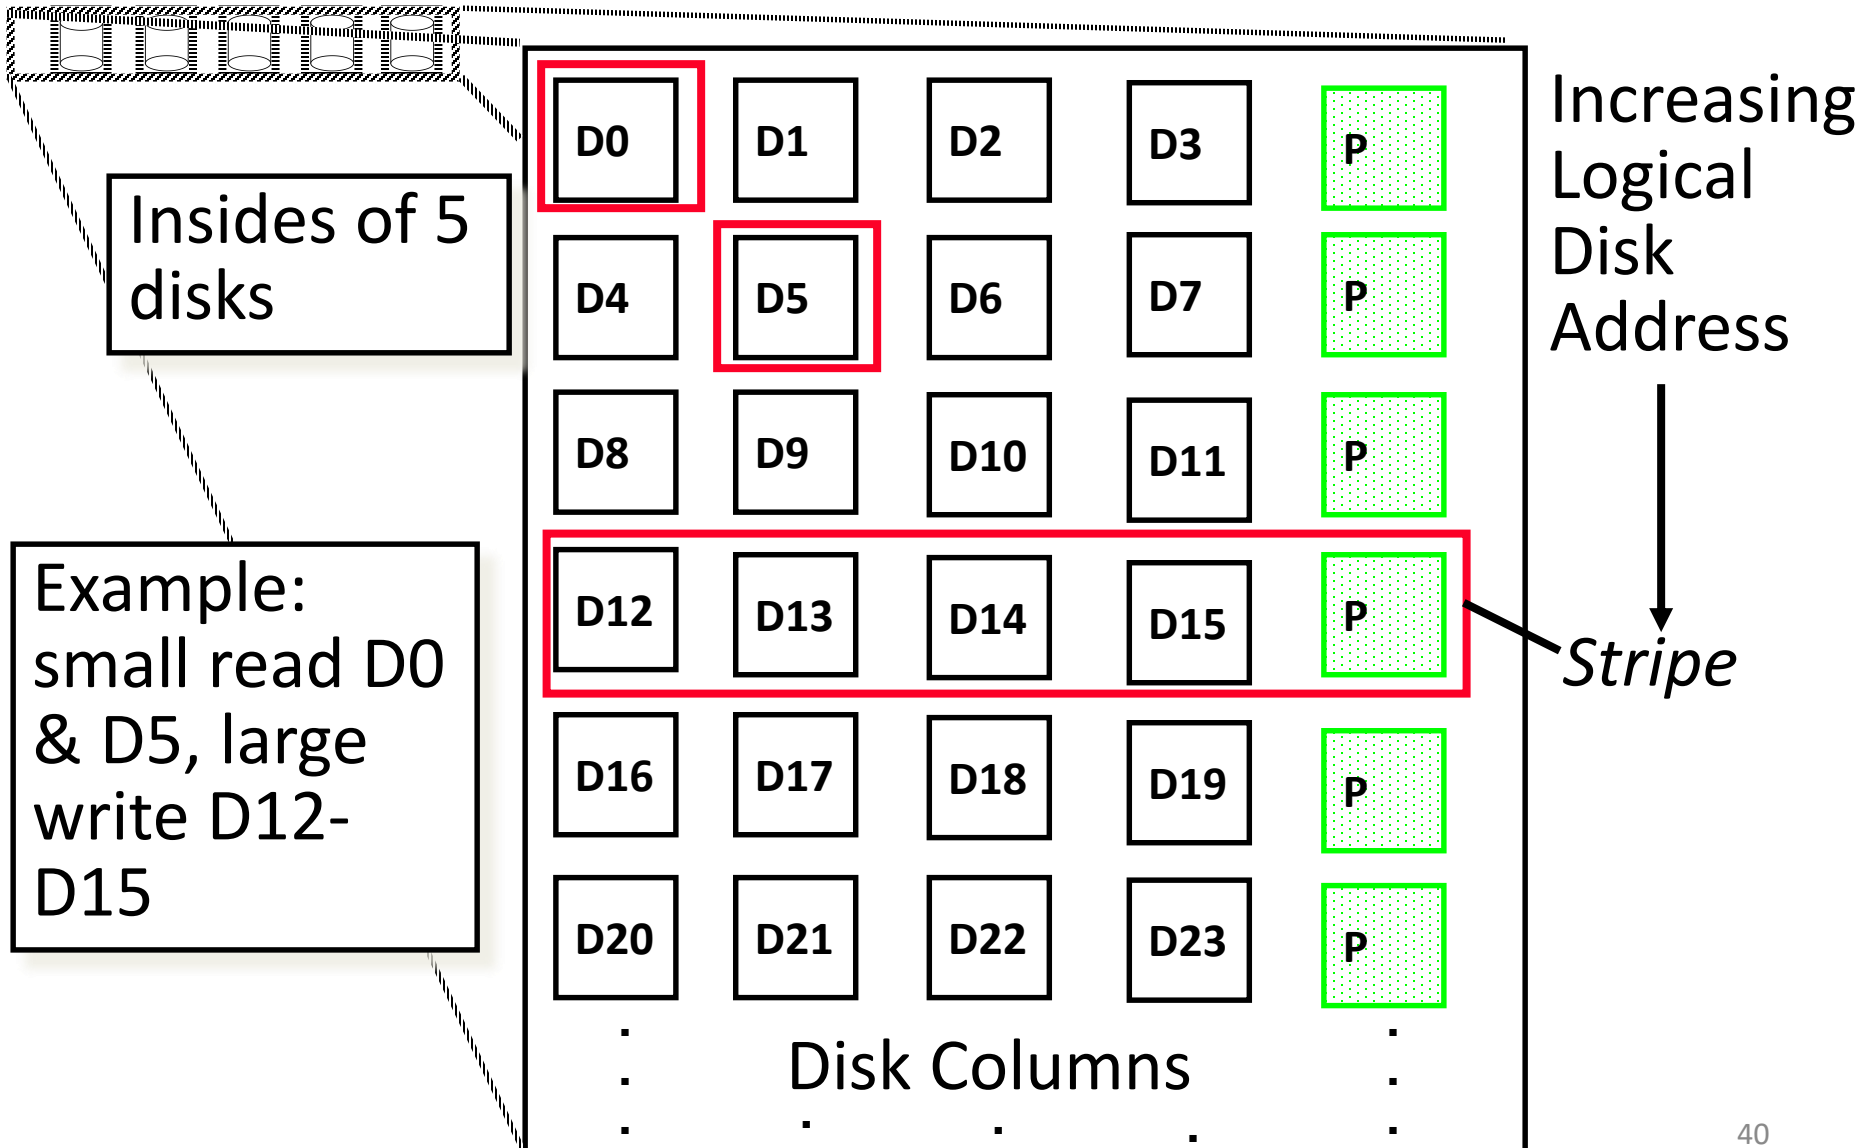


- Each disk is fully duplicated onto its “mirror(s)”
  - Very high availability can be achieved
- Bandwidth sacrifice on write:
  - Logical write = N physical writes
  - Reads may be optimized
- Most expensive solution: 100% capacity overhead
- RAID 10 (striped mirrors), RAID 01 (mirrored stripes):
  - Combinations of RAID 0 and 1.

# RAID 3: Parity Disk




# RAID 4: High I/O Rate Parity




# Inspiration for RAID 5

- RAID 4 works well for small reads
- Small writes (write to one disk):
  - Option 1: read other data disks, create new sum and write to Parity Disk
  - Option 2: since P has old sum, compare old data to new data, add the difference to P
- Small writes are limited by Parity Disk: Write to D0, D5 both also write to P disk



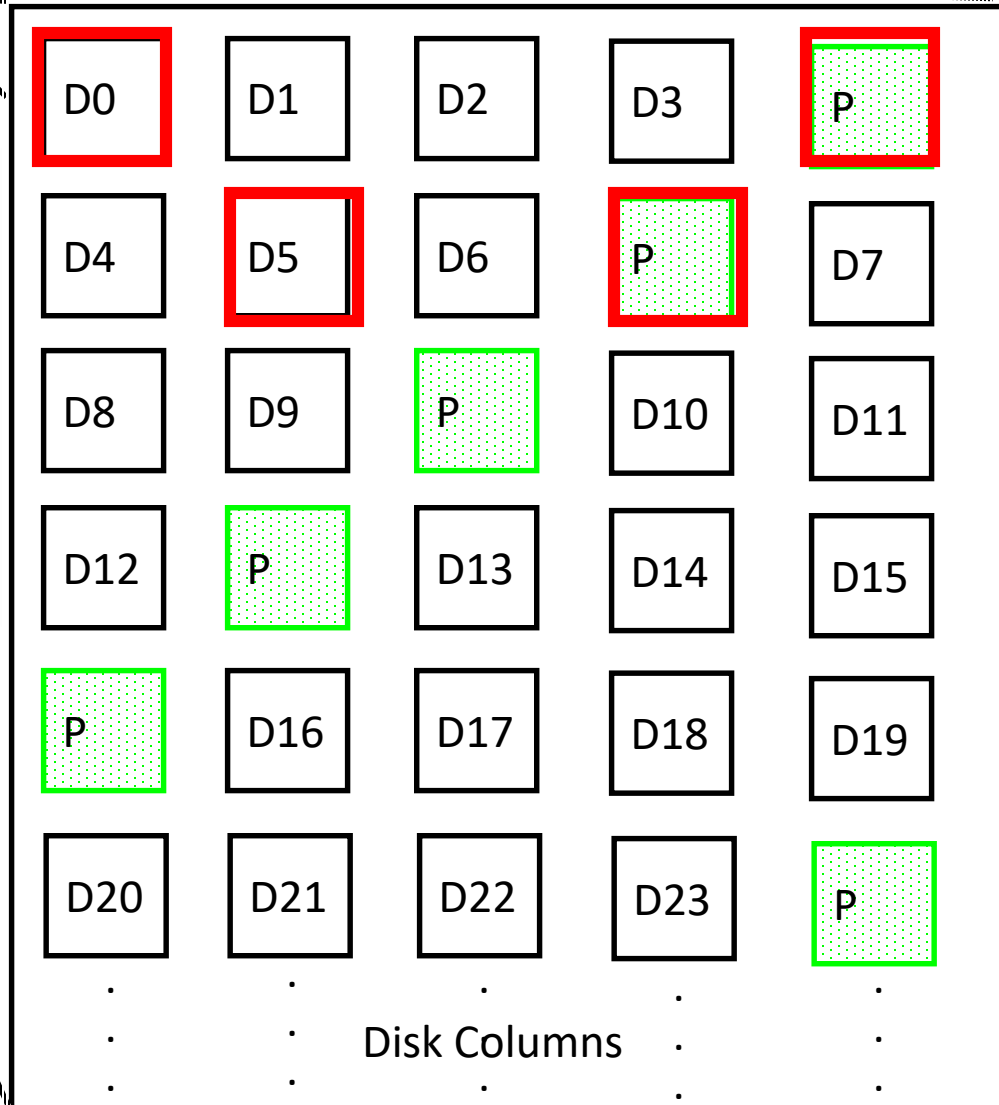


# RAID 5: High I/O Rate Interleaved Parity



Independent writes possible because of interleaved parity

Example:  
write to D0,  
D5 uses disks  
0, 1, 3, 4



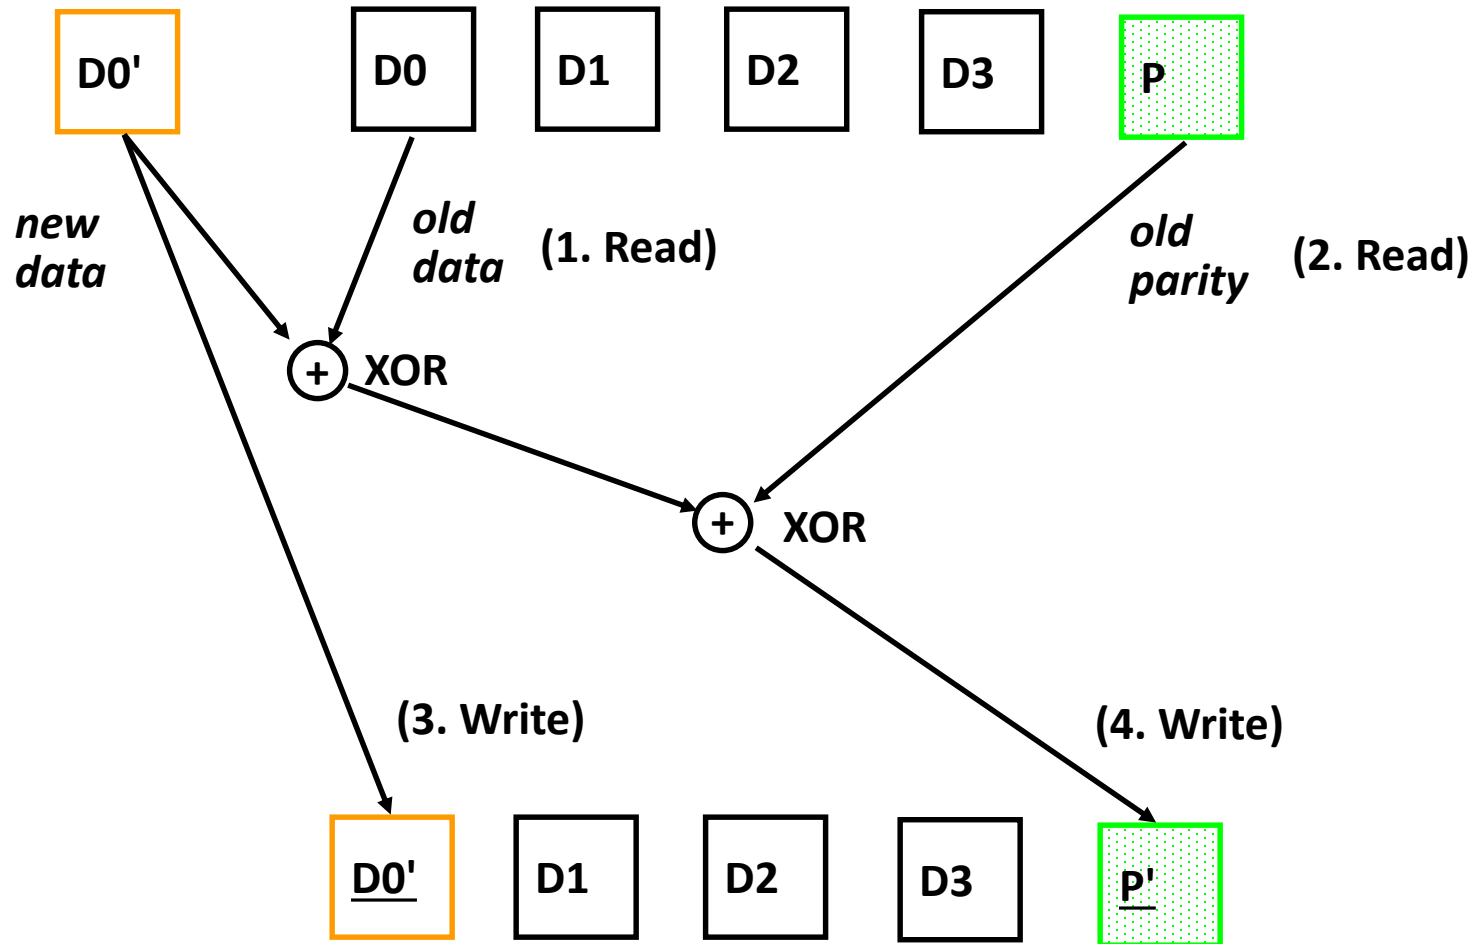
Increasing  
Logical  
Disk  
Addresses



# Problems of Disk Arrays: Small Writes

## RAID-5: *Small Write* Algorithm

1 Logical Write = 2 Physical Reads + 2 Physical Writes



# And, in Conclusion, ...

- Great Idea: Redundancy to Get Dependability
  - Spatial (extra hardware) and Temporal (retry if error)
- Reliability: MTTF & Annualized Failure Rate (AFR)
- Availability: % uptime
- Memory
  - Hamming ECC: correct single, detect double
- RAID
  - Interleaved data and parity