

2

**Solutions**

**2.1** `addi x5, x7, -5`  
`add x5, x5, x6`  
`[addi f, h, -5 (note, no subi) add f, f, g]`

**2.2** `f = g+h+i`

**2.3** `sub x30, x28, x29 // compute i-j`  
`slli x30, x30, 3 // multiply by 8 to convert the`  
`word offset to a byte offset`  
`ld x30, 0(x3) // load A[i-j]`  
`sd x30, 64(x11) // store in B[8]`

**2.4** `B[g] = A[f] + A[f+1]`  
`slli x30, x5, 3 // x30 = f*8`  
`add x30, x10, x30 // x30 = &A[f]`  
`slli x31, x6, 3 // x31 = g*8`  
`add x31, x11, x31 // x31 = &B[g]`  
`ld x5, 0(x30) // f = A[f]`  
`addi x12, x30, 8 // x12 = &A[f]+8 (i.e. &A[f+1])`  
`ld x30, 0(x12) // x30 = A[f+1]`  
`add x30, x30, x5 // x30 = A[f+1] + A[f]`  
`sd x30, 0(x31) // B[g] = x30 (i.e. A[f+1] + A[f])`

## 2.5

Little-Endian		Big-Endian	
Address	Data	Address	Data
12	ab	12	12
8	cd	8	ef
4	ef	4	cd
0	12	0	ab

## 2.6 2882400018

**2.7** `slli x28, x28, 3 // x28 = i*8`  
`ld x28, 0(x10) // x28 = A[i]`  
`slli x29, x29, 3 // x29 = j*8`  
`ld x29, 0(x11) // x29 = B[j]`  
`add x29, x28, x29 // Compute x29 = A[i] + B[j]`  
`sd x29, 64(x11) // Store result in B[8]`

**2.8**  $f = 2*(&A)$ 

```

addi x30, x10, 8    // x30 = &A[1]
addi x31, x10, 0    // x31 = &A
sd   x31, 0(x30)    // A[1] = &A
ld   x30, 0(x30)    // x30 = A[1] = &A
add  x5, x30, x31    // f = &A + &A = 2*(&A)

```

**2.9**

	type	opcode, funct3,7	rs1	rs2	rd	imm
addi x30,x10,8	I-type	0x13, 0x0, --	10	--	30	8
addi x31,x10,0	R-type	0x13, 0x0, --	10	--	31	0
sd x31,0(x30)	S-type	0x23, 0x3, --	31	30	--	0
ld x30,0(x30)	I-type	0x3, 0x3, --	30	--	30	0
add x5, x30, x31	R-type	0x33, 0x0, 0x0	30	31	5	--

**2.10****2.10.1** 0x5000000000000000**2.10.2** overflow**2.10.3** 0xB000000000000000**2.10.4** no overflow**2.10.5** 0xD000000000000000**2.10.6** overflow**2.11****2.11.1** There is an overflow if  $128 + x6 > 2^{63} - 1$ .In other words, if  $x6 > 2^{63} - 129$ .There is also an overflow if  $128 + x6 < -2^{63}$ .In other words, if  $x6 < -2^{63} - 128$  (which is impossible given the range of  $x6$ ).**2.11.2** There is an overflow if  $128 - x6 > 2^{63} - 1$ .In other words, if  $x6 < -2^{63} + 129$ .There is also an overflow if  $128 - x6 < -2^{63}$ .In other words, if  $x6 > 2^{63} + 128$  (which is impossible given the range of  $x6$ ).**2.11.3** There is an overflow if  $x6 - 128 > 2^{63} - 1$ .In other words, if  $x6 < 2^{63} + 127$  (which is impossible given the range of  $x6$ ).There is also an overflow if  $x6 - 128 < -2^{63}$ .In other words, if  $x6 < -2^{63} + 128$ .**2.12** R-type: add x1, x1, x1

**2.13** S-type: 0x25F3023 (0000 0010 0101 1111 0011 0000 0010 0011)

**2.14** R-type: sub x6, x7, x5 (0x40538333: 0100 0000 0101 0011 1000 0011 0011 0011)

**2.15** I-type: ld x3, 4(x27) (0x4DB183: 0000 0000 0100 1101 1011 0001 1000 0011)

## 2.16

**2.16.1** The opcode would expand from 7 bits to 9.

The rs1, rs2, and rd fields would increase from 5 bits to 7 bits.

**2.16.2** The opcode would expand from 7 bits to 12.

The rs1 and rd fields would increase from 5 bits to 7 bits. This change does not affect the imm field *per se*, but it might force the ISA designer to consider shortening the immediate field to avoid an increase in overall instruction size.

**2.16.3** \* Increasing the size of each bit field potentially makes each instruction longer, potentially increasing the code size overall.

\* However, increasing the number of registers could lead to less register spillage, which would reduce the total number of instructions, possibly reducing the code size overall.

## 2.17

**2.17.1** 0x1234567ababefef8

**2.17.2** 0x2345678123456780

**2.17.3** 0x545

**2.18** It can be done in eight RISC-V instructions:

```
addi x7, x0, 0x3f // Create bit mask for bits 16 to 11
slli x7, x7, 11   // Shift the masked bits
and  x28, x5, x7  // Apply the mask to x5
slli x7, x6, 15   // Shift the mask to cover bits 31
                  // to 26
xori x7, x7, -1   // This is a NOT operation
and  x6, x6, x7   // "Zero out" positions 31 to
                  // 26 of x6
slli x28, x28, 15 // Move selection from x5 into
                  // positions 31 to 26
or   x6, x6, x28  // Load bits 31 to 26 from x28
```

**2.19** xori x5, x6, -1

**2.20** `ld x6, 0(x17)`  
`slli x6, x6, 4`

**2.21** `x6 = 2`

**2.22**

**2.22.1** `[0x1ff00000, 0x200FFFFE]`

**2.22.2** `[0x1FFFF000, 0x20000ffe]`

**2.23**

**2.23.1** The UJ instruction format would be most appropriate because it would allow the maximum number of bits possible for the “loop” parameter, thereby maximizing the utility of the instruction.

**2.23.2** It can be done in three instructions:

```
loop:
    addi x29, x29, -1 // Subtract 1 from x29
    bgt x29, x0, loop // Continue if x29 not
                       negative
    addi x29, x29, 1 // Add back 1 that shouldn't
                     have been subtracted.
```

**2.24**

**2.24.1** The final value of `xs` is 20.

**2.24.2** `acc = 0;`  
`i = 10;`  
`while (i != 0) {`  
    `acc += 2;`  
    `i--;`  
`}`

**2.24.3** `4*N + 1` instructions.

**2.24.4** (Note: change condition `!=` to `>=` in the while loop)

```
acc = 0;
i = 10;
while (i >= 0) {
    acc += 2;
    i--;
}
```

**2.25** The C code can be implemented in RISC-V assembly as follows.

```

LOOPI:
    addi x7, x0, 0      // Init i = 0
    bge x7, x5, ENDI    // While i < a
    addi x30, x10, 0     // x30 = &D
    addi x29, x0, 0     // Init j = 0
LOOPJ:
    bge x29, x6, ENDJ    // While j < b
    add x31, x7, x29     // x31 = i+j
    sd x31, 0(x30)       // D[4*j] = x31
    addi x30, x30, 32    // x30 = &D[4*(j+1)]
    addi x29, x29, 1     // j++
    jal x0, LOOPJ
ENDJ:
    addi x7, x7, 1      // i++;
    jal x0, LOOPI
ENDI:

```

**2.26** The code requires 13 RISC-V instructions. When  $a = 10$  and  $b = 1$ , this results in 123 instructions being executed.

**2.27** // This C code corresponds most directly to the given assembly.

```

int i;
for (i = 0; i < 100; i++) {
    result += *MemArray;
    MemArray++;
}
return result;

```

// However, many people would write the code this way:

```

int i;
for (i = 0; i < 100; i++) {
    result += MemArray[i];
}
return result;

```

**2.28** The address of the last element of MemArray can be used to terminate the loop:

```
add x29, x10, 800      // x29 = &MemArray[101]
LOOP:
    ld    x7,    0(x10)
    add   x5,    x5, x7
    addi  x10,   x10, 8
    blt   x10,   x29, LOOP // Loop until MemArray points
                           // to one-past the last element
```

**2.29**

// IMPORTANT! Stack pointer must remain a multiple of 16!!!!

```
fib:
    beq   x10,   x0, done // If n==0, return 0
    addi  x5,    x0, 1
    beq   x10,   x5, done // If n==1, return 1
    addi  x2,    x2, -16 // Allocate 2 words of stack
                        // space
    sd    x1,    0(x2) // Save the return address
    sd    x10,   8(x2) // Save the current n
    addi  x10,   x10, -1 // x10 = n-1
    jal   x1,    fib // fib(n-1)
    ld    x5,    8(x2) // Load old n from the stack
    sd    x10,   8(x2) // Push fib(n-1) onto the stack
    addi  x10,   x5, -2 // x10 = n-2
    jal   x1,    fib // Call fib(n-2)
    ld    x5,    8(x2) // x5 = fib(n-1)
    add   x10,   x10, x5 // x10 = fib(n-1)+fib(n-2)
// Clean up:
    ld    x1,    0(x2) // Load saved return address
    addi  x2,    x2, 16 // Pop two words from the stack
done:
    jalr  x0,    x1
```

**2.30** [answers will vary]

**2.31**

```
// IMPORTANT! Stack pointer must remain a multiple of 16!!!
f:
    addi x2, x2, -16    // Allocate stack space for 2 words
    sd   x1, 0(x2)     // Save return address
    add  x5, x12, x13   // x5 = c+d
    sd   x5, 8(x2)     // Save c+d on the stack
    jal  x1, g          // Call x10 = g(a,b)
    ld   x11, 8(x2)    // Reload x11= c+d from the stack
    jal  x1, g          // Call x10 = g(g(a,b), c+d)
    ld   x1, 0(x2)     // Restore return address
    addi x2, x2, 16    // Restore stack pointer
    jalr x0, x1
```

**2.32** We can use the tail-call optimization for the second call to g, saving one instruction:

```
// IMPORTANT! Stack pointer must remain a multiple of 16!!!
f:
    addi x2, x2, -16    // Allocate stack space for 2 words
    sd   x1, 0(x2)     // Save return address
    add  x5, x12, x13   // x5 = c+d
    sd   x5, 8(x2)     // Save c+d on the stack
    jal  x1, g          // Call x10 = g(a,b)
    ld   x11, 8(x2)    // Reload x11 = c+d from the stack
    ld   x1, 0(x2)     // Restore return address
    addi x2, x2, 16    // Restore stack pointer
    jal  x0, g          // Call x10 = g(g(a,b), c+d)
```

**2.33** \*We have no idea what the contents of x10-x14 are, g can set them as it pleases.

\*We don't know what the precise contents of x8 and sp are; but we do know that they are identical to the contents when f was called.

\*Similarly, we don't know what the precise contents of x1 are; but, we do know that it is equal to the return address set by the "jal x1, f" instruction that invoked f.



**2.34**

```

a_to_i:
    addi    x28, x0, 10      # Just stores the constant 10
    addi    x29, x0, 0      # Stores the running total
    addi    x5, x0, 1       # Tracks whether input is positive
                             # or negative

    # Test for initial '+' or '-'
    lbu     x6, 0(x10)       # Load the first character
    addi    x7, x0, 45      # ASCII '-'
    bne     x6, x7, noneg
    addi    x5, x0, -1      # Set that input was negative
    addi    x10, x10, 1     # str++
    jal     x0, main_atoi_loop

noneg:
    addi    x7, x0, 43      # ASCII '+'
    bne     x6, x7, main_atoi_loop
    addi    x10, x10, 1     # str++

main_atoi_loop:
    lbu     x6, 0(x10)       # Load the next digit
    beq     x6, x0, done    # Make sure next char is a digit,
                             # or fail
    addi    x7, x0, 48      # ASCII '0'
    sub     x6, x6, x7
    blt     x6, x0, fail    # *str < '0'
    bge     x6, x28, fail    # *str >= '9'

    # Next char is a digit, so accumulate it into x29
    mul     x29, x29, x28    # x29 *= 10
    add     x29, x29, x6     # x29 += *str - '0'
    addi    x10, x10, 1     # str++
    jal     x0, main_atoi_loop

done:
    addi    x10, x29, 0     # Use x29 as output value
    mul     x10, x10, x5     # Multiply by sign
    jalr    x0, x1         # Return result

fail:
    addi    x10, x0, -1
    jalr    x0, x1

```

**2.35****2.35.1** 0x11**2.35.2** 0x88

```

2.36 lui    x10, 0x11223
      addi   x10, x10, 0x344
      slli   x10, x10, 32
      lui    x5, 0x55667
      addi   x5, x5, 0x788
      add    x10, x10, x5

```

**2.37**

```

setmax:
    try:
        lr.d  x5, (x10)           # Load-reserve *shvar
        bge   x5, x11, release    # Skip update if *shvar > x
        addi  x5, x11, 0
    release:
        sc.d  x7, x5, (x10)
        bne   x7, x0, try         # If store-conditional failed,
                                try again
        jalr  x0, x1

```

**2.38** When two processors A and B begin executing this loop at the same time, at most one of them will execute the store-conditional instruction successfully, while the other will be forced to retry the loop. If processor A's store-conditional succeeds initially, then B will re-enter the try block, and it will see the new value of shvar written by A when it finally succeeds. The hardware guarantees that both processors will eventually execute the code completely.

**2.39**

**2.39.1** No. The resulting machine would be slower overall.

Current CPU requires (num arithmetic \* 1 cycle) + (num load/store \* 10 cycles) + (num branch/jump \* 3 cycles) =  $500 * 1 + 300 * 10 + 100 * 3 = 3800$  cycles.

The new CPU requires  $(.75 * \text{num arithmetic} * 1 \text{ cycle}) + (\text{num load/store} * 10 \text{ cycles}) + (\text{num branch/jump} * 3 \text{ cycles}) = 375 * 1 + 300 * 10 + 100 * 3 = 3675$  cycles.

However, given that each of the new CPU's cycles is 10% longer than the original CPU's cycles, the new CPU's 3675 cycles will take as long as  $4042.5$  cycles on the original CPU.

**2.39.2** If we double the performance of arithmetic instructions by reducing their CPI to 0.5, then the CPU will run the reference program in  $(500 * .5) + (300 * 10) + 100 * 3 = 3550$  cycles. This represents a speedup of 1.07.

If we improve the performance of arithmetic instructions by a factor of 10 (reducing their CPI to 0.1), then the CPU will run the reference program in  $(500 * .1) + (300 * 10) + 100 * 3 = 3350$  cycles. This represents a speedup of 1.13.



